



Ecole d'ingénieurs et d'architectes de Fribourg
Hochschule für Technik und Architektur Freiburg

File Systems, LUKS, InitRamFS

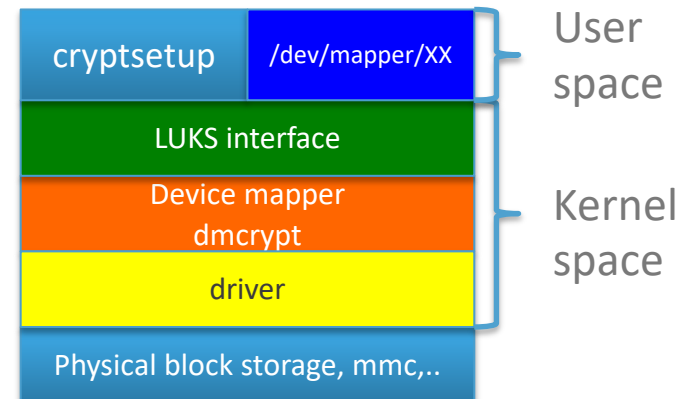
Youtube: <https://youtu.be/UoA6rVlbUJA>

References

- [1]: <Linux Kernel sources>/Documentation/filesystems
- [2]: http://www.tldp.org/HOWTO/html_single/SquashFS-HOWTO
- [3]: <http://squashfs.sourceforge.net>
- [4]:
tree.celinuxforum.org/CelfPubWiki/ELCEurope2008Presentations?action=AttachFile&do=get&target=squashfs-elce.pdf
- [5]: <http://superuser.com/questions/228657/which-linux-filesystem-works-best-with-ssd> //File for SSD card
- [6]: https://access.redhat.com/documentation/en-US/Red_Hat_Enterprise_Linux/6/html/Storage_Administration_Guide/index.html // very good site
- [7]: <https://code.google.com/p/cryptsetup/>
Power off embedded FS
- [8]: <http://stackoverflow.com/questions/14460091/embedded-file-system-and-power-off>
- [9]: https://elinux.org/images/0/02/Filesystem_Considerations_for_Embedded_Devices.pdf

LUKS, cryptsetup, dmccrypt

- See [7]: <https://code.google.com/p/cryptsetup/>
- <https://blog.tinned-software.net/automount-a-luks-encrypted-volume-on-system-start/>
- LUKS** (Linux Unified Key Setup) is the standard for **Linux hard disk encryption**.
- By providing a standard on-disk-format, it does not only facilitate compatibility among distributions, but also provides secure management of multiple user passwords.
- In contrast to existing solution, LUKS stores all necessary setup information **in the partition header, enabling the user to transport or migrate his data seamlessly**.
- LUKS – dmccrypt crypts an **entire partition**
- Luks features
 - compatibility via standardization,
 - secure against attacks,
 - support for multiple keys,
 - effective passphrase revocation,
 - free
- cryptsetup** is a utility used to configure **dmccrypt**
- cryptsetup uses the **/dev/random** and **/dev/urandom** node file

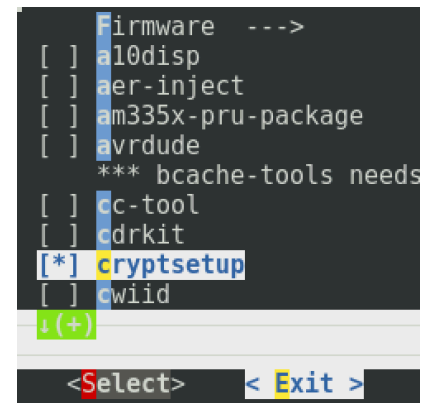


LUKS, cryptsetup, dmccrypt

- See : [7]: <https://code.google.com/p/cryptsetup/wiki/DMCrypt>
- **dmccrypt** (**Device-mapper**) crypts target and provides transparent encryption of block devices using the kernel crypto API (kernel configuration, Cryptographic API)
- Device-mapper is included in the Linux 2.6 and 3.x kernel that provides a generic way to create virtual layers of block devices. It is required by LVM2 (Logical Volume Management).
- The user can basically specify one of the symmetric ciphers, an encryption mode, a key (of any allowed size), an iv generation mode and then the user can create a new block device node file in **/dev/mapper**.
- All data written to this device will be encrypted and
- All data read from this device will be decrypted.

LUKS, cryptsetup, NanoPi, buildroot, kernel


- In order to **enable dmccrypt**, it is necessary to configure the **kernel**:
- `cd workspace/nano/buildroot`
`make linux-xconfig` or `make linux-menuconfig`
Go to: device driver → Multiple Devices drivers support (RAID and LVM) → Device mapper support → Crypt target support
- In order to use “cryptsetup”, it is required to add a new package in buildroot
`cd workspace/nano/buildroot`
`make menuconfig`
Go to: Target Packages → Hardware handling : choose `cryptsetup`



```
Firmware --->
[ ] al0disp
[ ] aer-inject
[ ] am335x-pru-package
[ ] avrdude
*** bcache-tools needs
[ ] cc-tool
[ ] cdrkit
[*] cryptsetup
[ ] cwiid
i (+)
<Select> <Exit>
```

LUKS with NanoPi

On the SD Card, create a third partition (with fdisk or parted)



Create LUKS partition

Initialize a LUKS partition, **be careful all data will be lost**. A passphrase generates the encryption key (--debug is optional)

On the NanoPi: **\$DEVICE** = /dev/mmcblk0p3

On PC: **\$DEVICE** = /dev/sdc3

Create a LUKS partition

```
# sudo cryptsetup --debug --pbkdf-memory 256 luksFormat $DEVICE
    Be careful, type yes in UPPERCASE
    (--pbkdf-memory limits the amount of memory)
```

Dump the header information of a LUKS device

```
# sudo cryptsetup luksDump /dev/mmcblk0p3
```

Create a mapping /dev/mapper/usrfs1 and ask the **passphrase**

```
# sudo cryptsetup --debug open --type luks $DEVICE usrfs1
```

Show the node file

```
# ls /dev/mapper/
brw----- 1 root  root    254,  0 Jan  1 15:36 usrfs1
```

SDCard

Usrfs,
partition 3

Create LUKS partition

Format the LUKS partition as ext4 partition

```
# sudo mkfs.ext4 /dev/mapper/usrfs1
```

Mount the LUKS partition to /mnt/usrfs

```
# sudo mkdir /mnt/usrfs
```

```
# sudo mount /dev/mapper/usrfs1 /mnt/usrfs
```

Work with the LUKS partition

```
# ls /mnt/usrfs
```

```
# copy files to /mnt/usrfs
```

Unmount the LUKS partition

```
# umount /dev/mapper/usrfs1
```

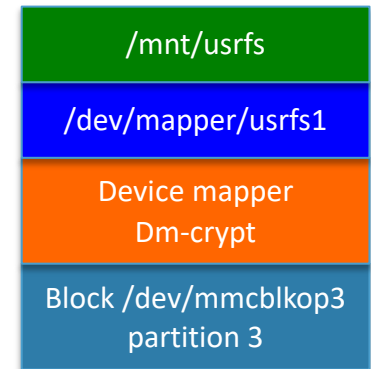
Removes the existing mapping `usrfs1` and wipes the key from kernel memory

```
# cryptsetup close usrfs1
```

dmsetup: low level logical volume management

```
# dmsetup info -C
```

```
# dmsetup remove -f usrfs1
```



Use LUKS partition

Create a mapping `/dev/mapper/usrfs1` and ask the **passphrase**

```
# sudo cryptsetup --debug open --type luks $DEVICE usrfs1
```

Mount the LUKS partition to `/mnt/usrfs`

```
# sudo mount /dev/mapper/usrfs1 /mnt/usrfs
```

Unmount the LUKS partition

```
# umount /dev/mapper/usrfs1
```

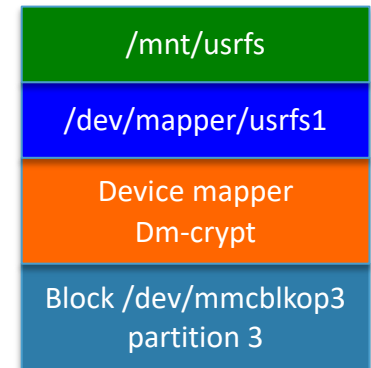
Removes the existing mapping `usrfs1` and wipes the key from kernel memory

```
# cryptsetup close usrfs1
```

It is possible to manage a luks partition with:

```
# dmsetup info -c
```

```
# dmsetup remove -f usrfs1
```



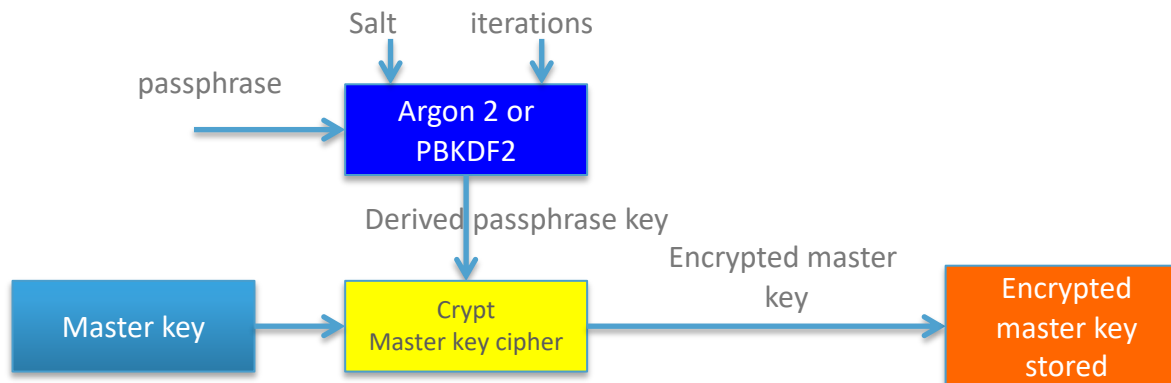
LUKS, key generation

- LUKS uses the **TKS1 template** in order to generate secure key.
- The key is derived from a passphrase
- LUKS supports multiple keys/passphrases
- TKS1 uses Argon2 or PBKDF2 (Password-Based Key Derivation Function 2) method in order to provide a better resistance against brute force attacks based on entropy weak user passphrase.
- TKS1 uses two level hierarchy of cryptographic keys to provide the ability to change passphrases
- See also: <http://clemens.endorphin.org/cryptography>

LUKS, key generation

The system initialization is straight forward:

- A master key is generated
- Passphrase, Salt, iterations and other values are inputs of the functions Argon2 or PBKDF2
- A derived passphrase key is computed by Argon2 or PBKDF2
- The master key is encrypted by the derived passphrase key.
- The encrypted master key, the iteration rate and the salt are stored



LUKS, key generation

Add a new passphrase to the LUKS partition

```
# cryptsetup luksAddKey /dev/mmcblk0p3
```

Dump the header information of a LUKS device

```
# cryptsetup luksDump /dev/mmcblk0p3
```

SDCard

Usrfs,
partition 3

```
Version:                1
Cipher name:             aes
Cipher mode:             xts-plain64
Hash spec:              sha1
Payload offset:         4096
MK bits:                256
MK digest:              6a ef 4e be 5d e5 90 80 48 fa a9 b0 21 cd cf be 9b cf 40 0e
MK salt:                d0 12 4d a2 52 80 72 fc 14 d2 f2 16 02 c5 e0 1d
                        9c 59 c4 fc 4e 9f 7b e7 be f6 b3 34 aa 09 ce 9c
MK iterations:          20125
UUID:                   d04071bc-d7e8-45d7-a950-a46c4e90d122
```

Crypt the master key

Key Slot 0: ENABLED

```
Iterations:             80000
Salt:                   bb e5 b8 ef 1d b4 03 5a f7 e5 1e 8e e0 70 d4 48
                        31 0c 31 52 b0 a4 2f 55 55 be 83 f2 ad c5 97 32
Key material offset:    8
AF stripes:             4000
```

Passphrase 1

Key Slot 1: ENABLED

```
Iterations:             81011
Salt:                   fd 21 0f d6 39 c4 1c 79 b5 2b ec 4d 43 dd 66 e0
                        ff 41 76 2f 39 59 e2 00 9a 2c 42 6b 22 10 16 47
Key material offset:    264
AF stripes:             4000
```

Passphrase 2

LUKS, key generation

Dump the encrypted master key of a LUKS device

```
# cryptsetup luksDump --dump-master-key /dev/mmcblk0p3
```

```
LUKS header information for /dev/mmcblk0p3
```

```
Cipher name:    aes
```

```
Cipher mode:    xts-plain64
```

```
Payload offset: 4096
```

```
UUID:          d04071bc-d7e8-45d7-a950-a46c4e90d122
```

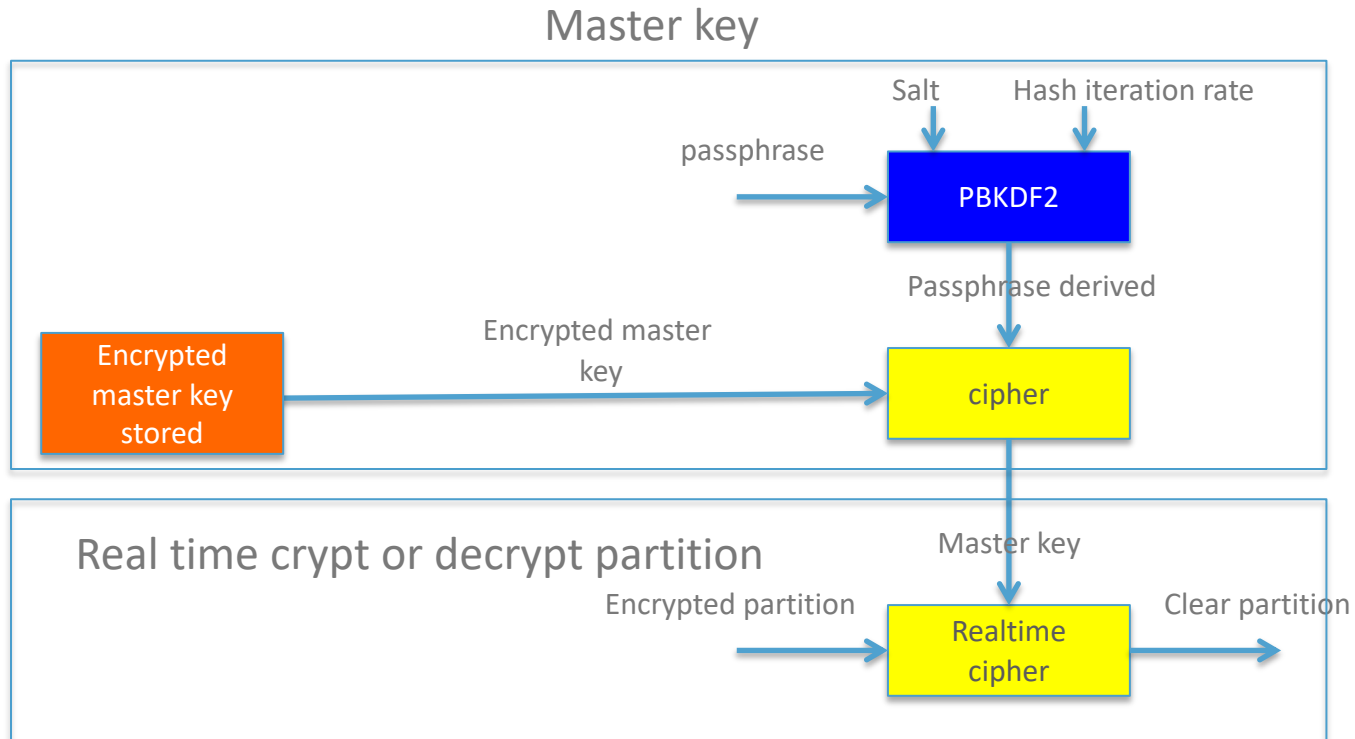
```
MK bits:       256
```

```
MK dump:       1e e2 d8 02 12 1a ce a4 74 66 20 3e 00 21 a7 c1  
                1b 92 88 76 d7 c1 d8 fd 1b 6e 42 fd ac 91 20 52
```

SDCard



LUKS, crypt partition



LUKS, cryptsetup, dmccrypt

- Check **/proc/crypto** which contains supported ciphers and modes (but note it contains only **currently loaded crypto API modules**).

```
Bash# cat /proc/crypto
```

```
...
name          : aes
driver        : aes-generic
module        : kernel
priority      : 100
refcnt        : 1
selftest      : passed
type          : cipher
blocksize     : 16
min keysize   : 16
max keysize   : 32
...
```

U-boot-Linux boot-without initramfs ^{1/2}

See u-boot course

Show boot.cmd file: `cat $HOME/workspace/nano/buildroot/board/friendlyarm/nanopi-neo-plus2/boot.cmd`

```
setenv bootargs console=ttyS0,115200n8 earlyprintk root=/dev/mmcblk0p2 rootwait
ext4load mmc 0 $kernel_addr_r Image
ext4load mmc 0 $fdt_addr_r nanopi-neo-plus2.dtb
booti $kernel_addr_r - $fdt_addr_r
```

Load Image

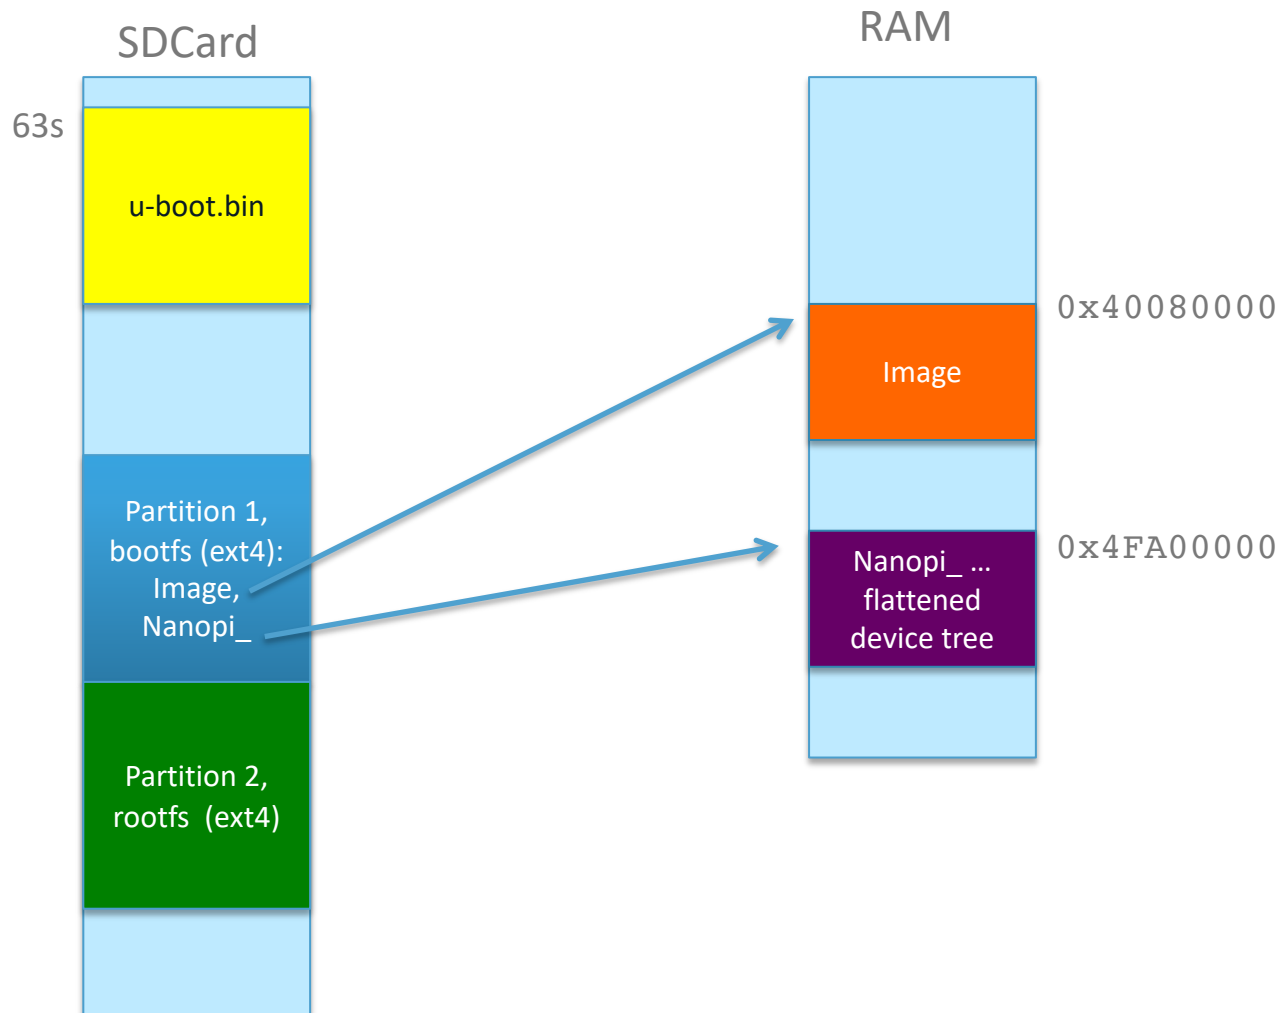
Load FDT

Start Linux

Linux kernel boot parameters

mmc 0: SDCard 1st partition

U-boot-Linux boot-**without** initramfs ^{2/2}



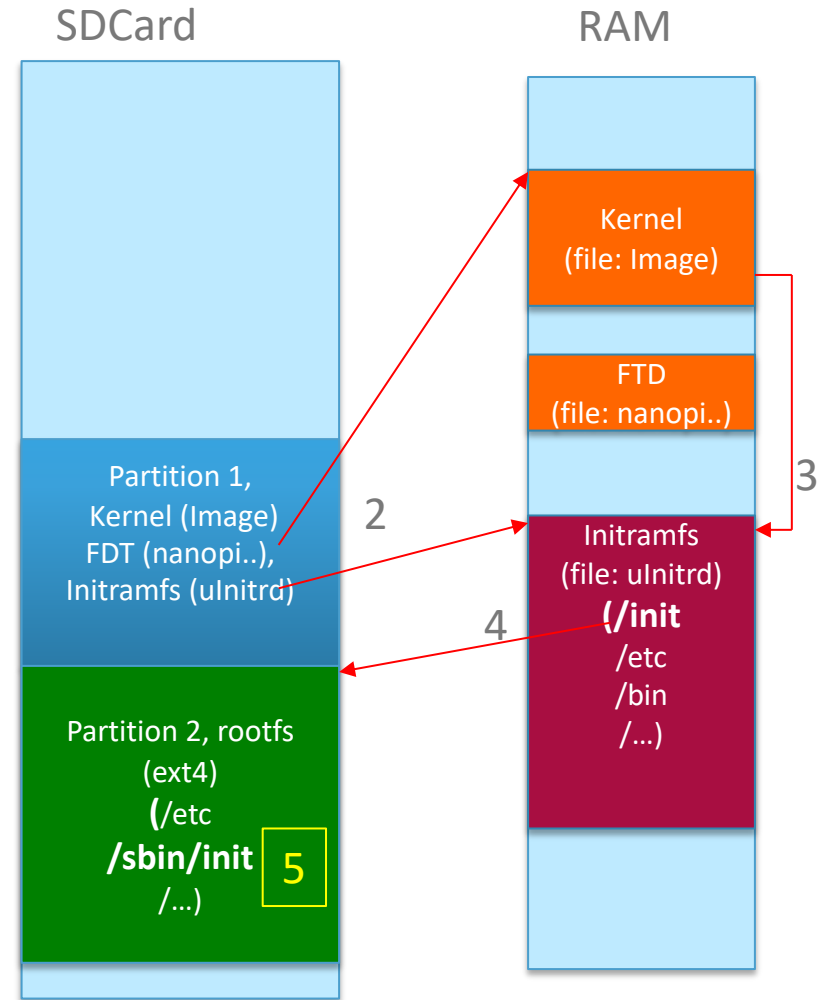
Initramfs

https://wiki.gentoo.org/wiki/Custom_Initramfs : good reference

- initramfs is a root filesystem that is **loaded at an early** stage of the boot process.
- It is the successor of initrd.
- It provides **early userspace commands** which lets the system do things that the kernel cannot easily do by itself during the boot process.
- Using initramfs is optional.
- Boot without initramfs:
 - By default, the kernel initializes hardware using built-in drivers, mounts the specified root partition, loads the rootfs and starts the init scripts
 - Init scripts can load additional modules and starts services until it eventually allows users to log in. This is a good default behavior and sufficient for many users.
- An initramfs is generally used for advanced requirements; for users who need to perform **certain tasks as early as possible, even before the rootfs is mounted.**

Boot-with initramfs

- 1) Kernel (Image), **initramfs (uInitrd)** and flattened device tree (Sun50i...) files are located in the partition 1 of the SDCard
- 2) Kernel, initramfs, Sun50i.. are copied to the RAM
- 3) Kernel mounts initramfs (uInitrd file)
- 4) Kernel executes **init** script stored in **initramfs**. This init script can execute early different commands
- 5) Init script executes the command `switch_root`, which switches to the standard rootfs located in the partition 2 and executes the `/sbin/init` command



U-boot-Linux boot-with initramfs 1/2

Show boot.cmd file: `cat $HOME/workspace/nano/buildroot/board/friendlyarm/nanopi-neo-plus2/boot.cmd`

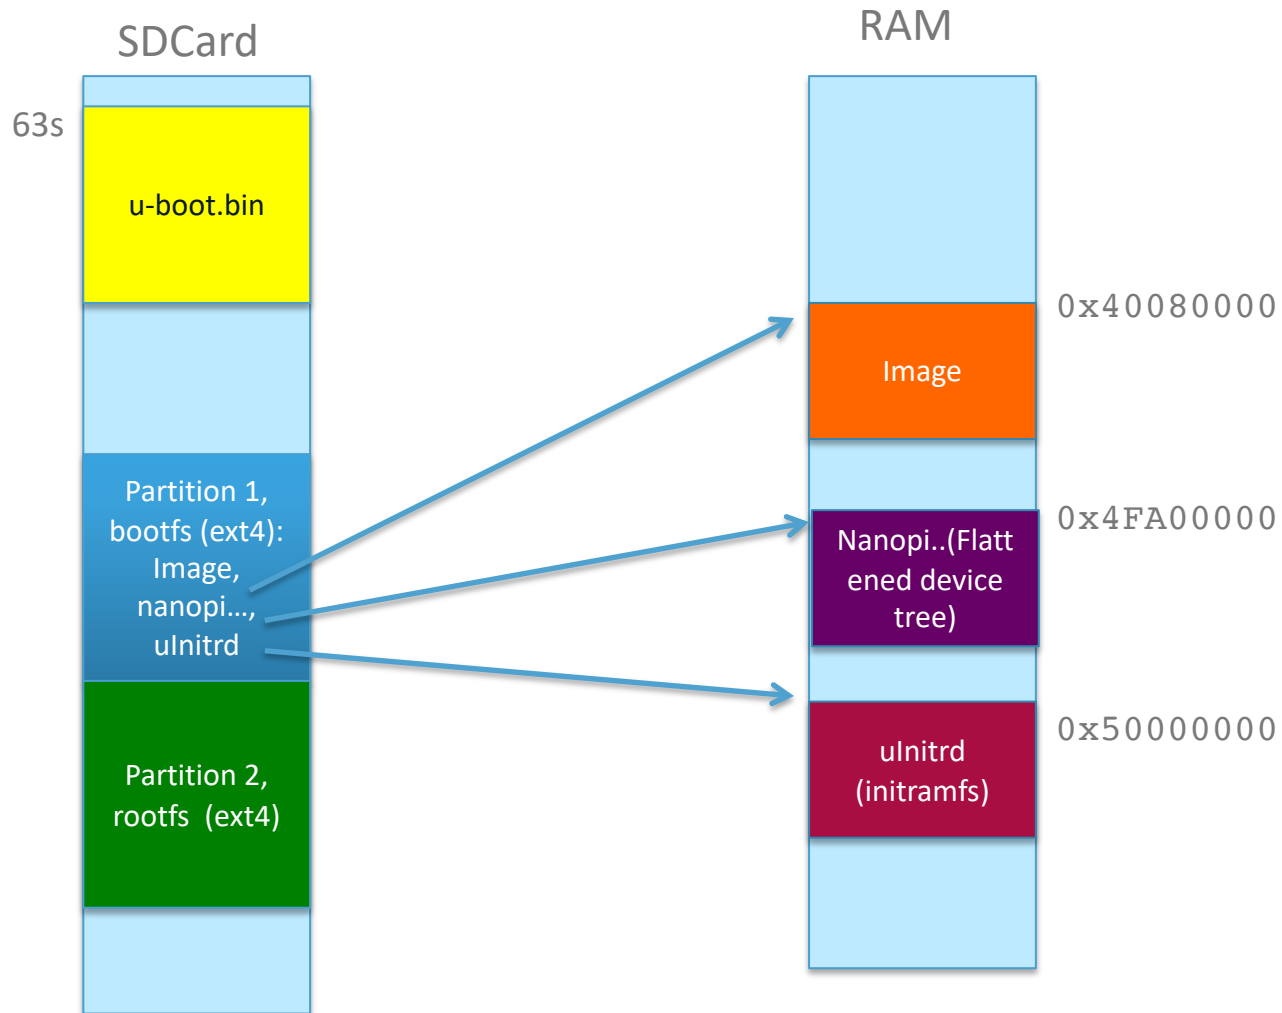
```
setenv bootargs console=ttyS0,115200n8 earlyprintk root=/dev/mmcblk0p2 rootwait
ext4load mmc 0 $kernel_addr_r Image
ext4load mmc 0 $fdt_addr_r nanopi-neo-plus2.dtb
```

```
ext4load mmc 0 0x50000000 uInitrd           // Load initramfs
```

```
booti $kernel_addr_r 0x50000000 $fdt_addr_r
```

initramfs address

U-boot-Linux boot-with initramfs ^{2/2}




How to build Initramfs

On PC, NanoPi rootfs is in this directory: `$HOME/workspace/nano/buildroot/output/target/`

Principle to build an initramfs:

1. to copy the right files into a directory (`$HOME/ramfs`),
2. to copy these files in a cpio archive file,
3. to compress this file
4. To add the uboot header




Initramfs: kernel configuration

Kernel configuration:

`CONFIG_BLK_DEV_INITRD=y`

General setup ---> [*] Initial RAM filesystem and RAM disk (initramfs/initrd) support



Automount a devtmpfs and initiate the /dev:

Device Drivers → Generic Drivers options → Maintain a devtmpfs filesystem to mount at /dev
→ Automount a devtmpfs

(Generally this option is not used)


Embedding the initramfs into the kernel:

`CONFIG_INITRAMFS_SOURCE="/usr/src/initramfs"`

General setup ---> [*] Initial RAM filesystem and RAM disk (initramfs/initrd) support

Initramfs manual generation ^{1/11}

- initramfs is a **cpio archive file**. It can be generated automatically with genkernel or dracut commands.
- initramfs can be manually generated.
- An *initramfs* contains at *least* one file called **/init**.
- kernel function `start_kernel()` (<Linux sources>/init.main.c) searches and executes the **/init program or script**



/init script 2/11

```
#!/bin/busybox sh
# Init script in the initRamFS
mount -t proc none /proc
mount -t sysfs none /sys
```



Mount the /proc and /sys pseudo-filesystem

```
mount -t ext4 /dev/mmcblk0p2 /newroot
```

Mount the rootfs (2nd partition) to /newroot

```
mount -n -t devtmpfs devtmpfs /newroot/dev
```

Populate /newroot/dev

```
exec switch_root /newroot /sbin/init
```

Switch to the rootfs on partition 2 and execute the /sbin/init command

`man switch_root`: switch to another filesystem as the root of the mount tree

Shared library dependency 3/11

- Generally programs are dynamically linked (other possibility: statically linked)
- A dynamically program must have all necessary libraries.
- Example (on PC) for the `/bin/ls` program:

```
ldd ls
linux-vdso.so.1 => (0x00007fff46198000)
libselinux.so.1 => /lib64/libselinux.so.1 (0x00000032fca00000)
libcap.so.2 => /lib64/libcap.so.2 (0x0000003dd9c00000)
libacl.so.1 => /lib64/libacl.so.1 (0x0000003dd9800000)
libc.so.6 => /lib64/libc.so.6 (0x0000003dc0800000)
libdl.so.2 => /lib64/libdl.so.2 (0x0000003dc0c00000)
libpcre.so.1 => /lib64/libpcre.so.1 (0x00000032fc600000)
libpthread.so.0 => /lib64/libpthread.so.0 (0x0000003dc1000000)
/lib64/ld-linux-x86-64.so.2 (0x0000003dc0400000)
libattr.so.1 => /lib64/libattr.so.1 (0x0000003dd7800000)
```

- `ls` program needs the `linux-vdso.so.1`, `libselinux.so.1`, ... libraries
- These libraries are in the `/lib` or `/lib64` directories.
- It is possible to use `strings` command in order to find the library dependency

```
strings ls | grep lib
```

Shared library dependency 4/11

- strace command is another possibility to find dynamic libraries used by a program.
- strace shows used libraries and the path where these libraries are
- Example with the program cryptsetup on NanoPi

```
# strace -f cryptsetup luksFormat /dev/mmcblk1p3
openat(AT_FDCWD, "/lib64/libm.so.6", O_RDONLY|O_CLOEXEC) = 3
openat(AT_FDCWD, "/usr/lib64/libcryptsetup.so.12", O_RDONLY|O_CLOEXEC) = 3
openat(AT_FDCWD, "/usr/lib64/libpopt.so.0", O_RDONLY|O_CLOEXEC) = 3
openat(AT_FDCWD, "/lib64/libuuid.so.1", O_RDONLY|O_CLOEXEC) = 3
openat(AT_FDCWD, "/lib64/libblkid.so.1", O_RDONLY|O_CLOEXEC) = 3
openat(AT_FDCWD, "/lib64/libpthread.so.0", O_RDONLY|O_CLOEXEC) = 3
openat(AT_FDCWD, "/lib64/libc.so.6", O_RDONLY|O_CLOEXEC) = 3
openat(AT_FDCWD, "/usr/lib64/libdevmapper.so.1.02", O_RDONLY|O_CLOEXEC) = 3
openat(AT_FDCWD, "/usr/lib64/libssl.so.1.1", O_RDONLY|O_CLOEXEC) = 3
openat(AT_FDCWD, "/usr/lib64/libcrypto.so.1.1", O_RDONLY|O_CLOEXEC) = 3
openat(AT_FDCWD, "/usr/lib64/libjson-c.so.4", O_RDONLY|O_CLOEXEC) = 3
openat(AT_FDCWD, "/lib64/libdl.so.2", O_RDONLY|O_CLOEXEC) = 3
openat(AT_FDCWD, "/lib64/libatomic.so.1", O_RDONLY|O_CLOEXEC) = 3
```

Shared library dependency 5/11

- In order to limit the shared libraries, busybox program can be used (see course nanopi.pdf)
- On PC, the directory `~workspace/nano/buildroot/output/target` contains an image of the rootfs

```
Linux$ cd ~/workspace/nano/buildroot/output/target
Linux$ cd bin
Linux$ strings busybox | grep lib
      libc.so.6
      libresolv.so.2
      /lib/ld-linux-aarch64.so.1
```

Important: busybox needs these three libraries:

```
libc.so.6
libresolv.so.2
/lib/ld-linux-aarch64.so.1
```

It is necessary to know in which directories are these three libraries

Shared library dependency 5/11

- strace shows used libraries and the path where these libraries are

- Example: ls program on NanoPi (ls is a symbol link to busybox)

```
strace -f ls
```

```
openat(AT_FDCWD, "/lib64/libresolv.so.2", O_RDONLY|O_CLOEXEC) = 3
```

```
openat(AT_FDCWD, "/lib64/libc.so.6", O_RDONLY|O_CLOEXEC) = 3
```

We can see that ls (link to busybox) uses:

- /lib64/libresolv.so.2 and
- /lib64/libc.so.6 libraries

Shared library dependency 5/11

```
Linux$ cd ~/workspace/nano/buildroot/output/target
```

libc library

```
Linux$ cd lib64
```

```
-rwxr-xr-x. 1 root root 1414752 Nov  4 09:12 libc-2.30.so  
lrwxrwxrwx. 1 root root      12 Sep 14 16:23 libc.so.6 -> libc-2.30.so
```

libresolv library

```
Linux$ cd lib64
```

```
-rwxr-xr-x      1 root root 80392 Nov  4 2019 libresolv-2.30.so  
lrwxrwxrwx      1 root root   17 Sep 14 2019 libresolv.so.2 -> libresolv-2.30.so
```

Shared library dependency 5/11

```
Linux$ cd ~/workspace/nano/buildroot/output/target
```

ld-linux-aarch64 library

```
Linux$ cd lib64
```

```
lrwxrwxrwx. 1 root root 19 Sep 14 16:23 ld-linux-aarch64.so.1 -> ../lib64/ld-2.30.so
-rwxr-xr-x. 1 root root 159384 Nov  4 09:12 ld-2.30.so
```

```
Linux$ cd lib
```

```
-rwxr-xr-x. 1 root root 159384 Nov  4 09:12 ld-2.30.so
lrwxrwxrwx. 1 root root      19 Sep 14 16:23 ld-linux-aarch64.so.1 ->
../lib64/ld-2.30.so
```

Initramfs summary ^{7/11}

The initramfs contains these files and directories:

```
/bin
/bin/mknod
/bin/mount
/bin/ln
/bin/sleep
/bin/umount
/bin/ls
/bin/mkdir
/bin/sh
/bin/busybox
/sbin
/sbin/switch_root
```

Busybox and
symbolic links

```
/proc

/lib64/libc.so.6
/lib64/libc-2.30.so
/lib64/libresolv.so.2
/lib64/libresolv-2.30.so
/lib64/ld-linux-aarch64.so.1
/lib64/ld-2.30.so
/lib/ld-linux-aarch64.so.1
/lib/ld-2.30.so
```

Shared libraries
And symbolic links

```
/dev
/dev/null
/dev/console
/dev/tty
/dev/random
/dev/urandom
/dev/ttyS0
/dev/ttyS1
/dev/ttyS2
/dev/ttyS3
/dev/mmcblk0p
/dev/mmcblk0p1
/dev/mmcblk0p2
/dev/mmcblk0p3
/dev/mmcblk0p4
./etc
./home
```

Nodes files

```
./init    /init script (p 23)
./sys
./newroot
./root
```


Build initramfs 8/11

This script builds the initramfs in the directory **ROOTFSLOC=ramfs**

```
#!/bin/bash
ROOTFSLOC=ramfs
cd $HOME
mkdir $ROOTFSLOC
mkdir -p $ROOTFSLOC/{bin,dev,etc,home,lib,lib64,newroot,proc,root,sbin,sys}

cd $ROOTFSLOC/dev
sudo mknod null      c 1 3
sudo mknod tty       c 5 0
sudo mknod console  c 5 1
sudo mknod random    c 1 8
sudo mknod urandom   c 1 9

sudo mknod mmcblk0p  b 179 0
sudo mknod mmcblk0p1 b 179 1
sudo mknod mmcblk0p2 b 179 2
sudo mknod mmcblk0p3 b 179 3
sudo mknod mmcblk0p4 b 179 4

sudo mknod ttyS0     c 4 64
sudo mknod ttyS1     c 4 65
sudo mknod ttyS2     c 4 66
sudo mknod ttyS3     c 4 67
```

Nodes files

Build initramfs 9/11

```
cd ../bin
cp ~/workspace/nano/buildroot/output/target/bin/busybox .

ln -s busybox ls
ln -s busybox mkdir
ln -s busybox ln
ln -s busybox mknod
ln -s busybox mount
ln -s busybox umount
ln -s busybox sh
ln -s busybox sleep
ln -s busybox dmesg
cp ~/workspace/nano/buildroot/output/target/usr/bin/strace .
```

/bin
Busybox and
symbolic links
strace

```
cd ../sbin
ln -s ../bin/busybox switch_root
```

/sbin

Build initramfs 10/11

```
cd ../lib64
cp ~/workspace/nano/buildroot/output/target/lib64/ld-2.30.so .
cp ~/workspace/nano/buildroot/output/target/lib64/libresolv-2.30.so .
cp ~/workspace/nano/buildroot/output/target/lib64/libc-2.30.so .
ln -s libresolv-2.30.so libresolv.so.2
ln -s libc-2.30.so libc.so.6
ln -s ../lib64/ld-2.30.so ld-linux-aarch64.so.1

cd ../lib
cp ~/workspace/nano/buildroot/output/target/lib64/ld-2.30.so .
ln -s ../lib64/ld-2.30.so ld-linux-aarch64.so.1
```

Shared libraries

Build initramfs 10/11

```
cd ..
cat > init << endofinput
#!/bin/busybox sh

mount -t proc none /proc
mount -t sysfs none /sys
mount -t ext4 /dev/mmcblk0p2 /newroot
mount -n -t devtmpfs devtmpfs /newroot/dev

#exec sh
exec switch_root /newroot /sbin/init

endofinput
#####
chmod 755 init
```

/init

```
cd ..
sudo chown -R 0:0 $ROOTFSLOC
```

Change owner to root:root (0:0)

Build the initramfs cpio archive 11/11

```
cd $ROOTFSLOC
```

```
find . | cpio --quiet -o -H newc > ../Initrd
```

It is mandatory to be in the
ROOTFSLOC directory

```
cd ..
```

```
gzip -9 -c Initrd > Initrd.gz
```

Compress the initramfs

```
mkimage -A arm -T ramdisk -C none -d Initrd.gz uInitrd
```

Add the u-boot header