



Ecole d'ingénieurs et d'architectes de Fribourg
Hochschule für Technik und Architektur Freiburg


3: Compile Kernel

Software attacks

Buffer overflow:

Example:


```
void main () {  
    char buffer [4];  
    strcpy (&buffer[0], "123456");  
}
```



Software attacks

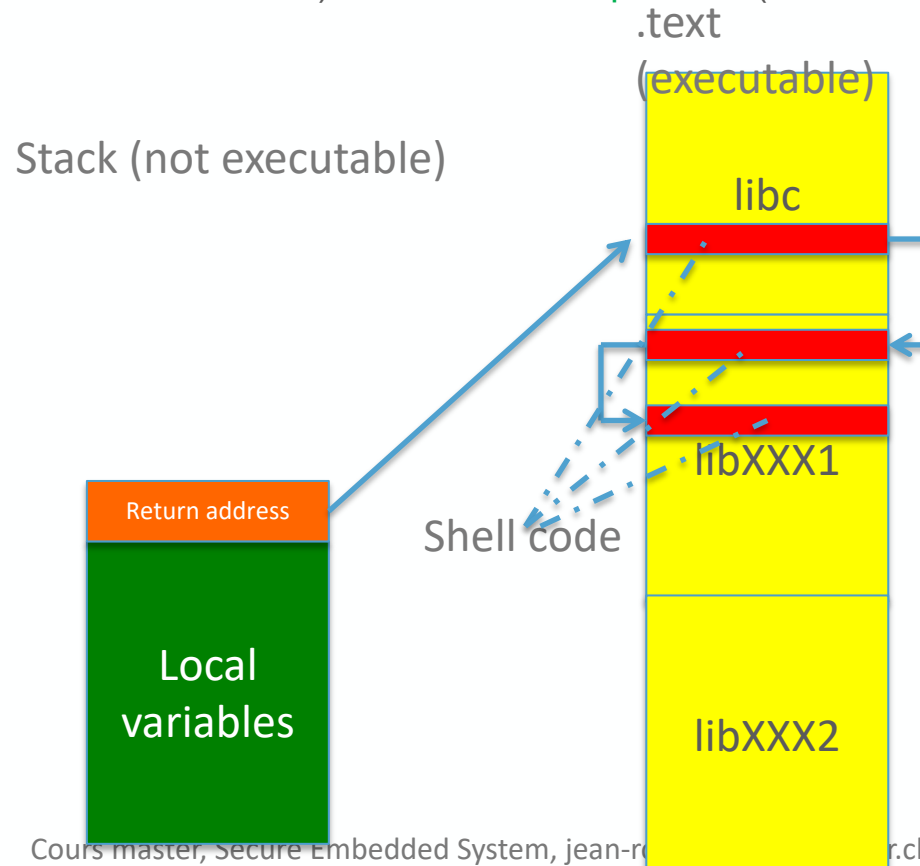
- A buffer overflow attack can insert and executes a shell code in an executable stack.

Stack (executable)




Software attacks

- Currently, the stack memory is no longer executable
- However, a technique called **ret2libc** could be used to bypass non executable memory. The main idea is to execute code in an executable memory like `libc()` or other libraries.
- Another technique called **ROP**, or Return-Oriented Programming allows also to bypass this protection. The main idea is to execute code in the program itself
- **ASLR** (Address Space Layout Randomization) `randomize_va_space` is used in order **avoid the ret2lib** (because stack and head addresses change)
- The **PIE** (Position Independent Executable) **avoids the ROP problem** (because code addresses change)



Linux spaces

- Generally users use the User space
- `dmesg`, `/dev/mem`, `/proc`, `/sys`, modules and devices drivers allow data exchange between User and Kernel spaces



Linux kernel

Linux kernel is in this directory:

```
~/workspace/nano/buildroot/output/build/linux-xx/
```

This directory has these main sub-directories:

arch	Hardware dependent code
block	Generic functions for the block devices
crypto	Cryptographic algo. used in the kernel
Documentation	Documentation about the kernel
drivers	All drivers known by the kernel
fs	All filesystem know by the kernel
include	kernel include files
init	Init code (function start_kernel)
ipc	Interprocess communication
kernel	Kernel code, scheduler, mutex, ...
lib	different libraries used by the kernel

Linux kernel

mm	Memory management
net	Different protocols, IPv4, IPv6, bluetooth, ...
samples	Different examples, kobject, kfifo, ...
security	Encrypted keys, SELinux, ...
sound	Sound support for Linux kernel
virt	Kernel-based virtual machine

This directory has these main files

vmlinux	Linux kernel, ELF format, ARM aarch64
.config	Linux kernel configuration
.config.old	Old Linux kernel configuration
Kconfig	Configuration for the make linux-xconfig
Makefile	makefile

Boot Linux kernel

During the startup phase, Linux executes first this file:

```
arch/arm64/kernel/head.S           // hardware dependent code
```

Next this function is called:

```
start_kernel( ) (init/main.c)  // hardware independent code
```

This function initializes different things:

- Boot CPU
- Memory management
- SMP, Symmetric Multi-processing
- Parse the boot parameters
- Interruptions
- Scheduler
- Timers
- Initial Ram Disk
- Key
- Security (SELinux, ...)

make friendlyarm_nanopi_neo_plus2_defconfig

Important: `cd ~/workspace/nano/buildroot`

During buildroot installation the command

`make friendlyarm_nanopi_neo_plus2_defconfig` was executed

The `make friendlyarm_nanopi_neo_plus2_defconfig` command initializes the `.config` file and indicates which processor, card, hardware, ... is used.

`.config` file:

```
CONFIG_CC_IS_GCC=y
CONFIG_GCC_VERSION=80201
CONFIG_CLANG_VERSION=0
CONFIG_CC_HAS_ASM_GOTO=y
CONFIG_CC_HAS_WARN_MAYBE_UNINITIALIZED=y
CONFIG_IRQ_WORK=y
CONFIG_BUILDTIME_EXTABLE_SORT=y
CONFIG_THREAD_INFO_IN_TASK=y
```

`#`

`# General setup`

`#`

```
CONFIG_INIT_ENV_ARG_LIMIT=32
# CONFIG_COMPILE_TEST is not set
CONFIG_LOCALVERSION=""
```

Compile linux

Compile

```
cd ~/workspace/nano/buildroot
```

```
make linux-menuconfig or // to configure the linux kernel
make linux-xconfig       // to configure the linux kernel
make linux-rebuild       // to compile only the linux kernel
```

Other commands

```
make clean // Be careful: delete output
           // directory and the .o files
           // → compile all files

make help // Different commands
make nano_neno_plus2_defconfig // Only the first time
make V=1 // print commands
```

make linux-xconfig

Important: `cd ~/workspace/nano/buildroot`

`make linux-xconfig` or `make linux-menuconfig` configures the kernel

`make linux-xconfig` (remark: install qt and qt-devel)

Make linux-xconfig

Kernel configuration is organized in different categories

- General setup
- Enable loadable module
- Enable the block layer
- ...

Patch physical to virtual translations at runtime

- ▶ General setup
- ▶ Enable loadable module support
- ▶ Enable the block layer
- ▶ System Type
- ▶ Bus support
- ▶ Kernel Features
- ▶ Boot options
- ▶ CPU Power Management
- ▶ Floating point emulation
- ▶ Userspace binary formats
- ▶ Power management options
- ▶ Networking support
- ▶ Device Drivers
- ▶ Firmware Drivers
- ▶ File systems
- ▶ Kernel hacking
- ▶ Security options
- ▶ Cryptographic API
- ▶ Library routines
- ▶ Virtualization

Disable Kernel .config support (CONFIG_IKCONFIG)

This option enables the complete Linux kernel ".config" file contents to be saved in the kernel. Kernel options can be read with this command: `cat /prog/config.gz`

Hardening Linux: Not include the configuration file inside the kernel

make linux-xconfig: General Setup → [] Kernel .config support

Enable `-fstack-protector` buffer overflow detection (CC_STACKPROTECTOR)

gcc `-fstack-protector-all` option adds extra code to check for buffer overflows, such as stack smashing attacks. This is done by adding a guard variable to functions.

Hardening Linux:

make linux-xconfig: General architecture-dependent options →

[*] Stack Protector buffer overflow detection

[*] Strong Stack Protector

Enabling this option provides some level of protection against stack based buffer overflows within the Linux kernel memory (not the user processes). If detected, the kernel will die with a kernel panic.

Enable -fstack-protector buffer overflow detection (CC_STACKPROTECTOR)

Example: file.c has an buffer overflow error, the `-fstack-protector-all` detects that

```
unsigned char localBufferOverflow (void) {  
    unsigned char val    [16];  
    int          i;  
    for (i=0; i< 18; i++)  
        val[i] = 0;  
    return val[0];  
}
```

```
# gcc -Wall -fstack-protector-all -o file file.C  
#./file  
*** stack smashing detected ***: ./file terminated  
===== Backtrace: =====  
/lib/libc.so.6(__fortify_fail+0x45)[0x4de0eb85]  
/lib/libc.so.6[0x4de0eb3a]
```

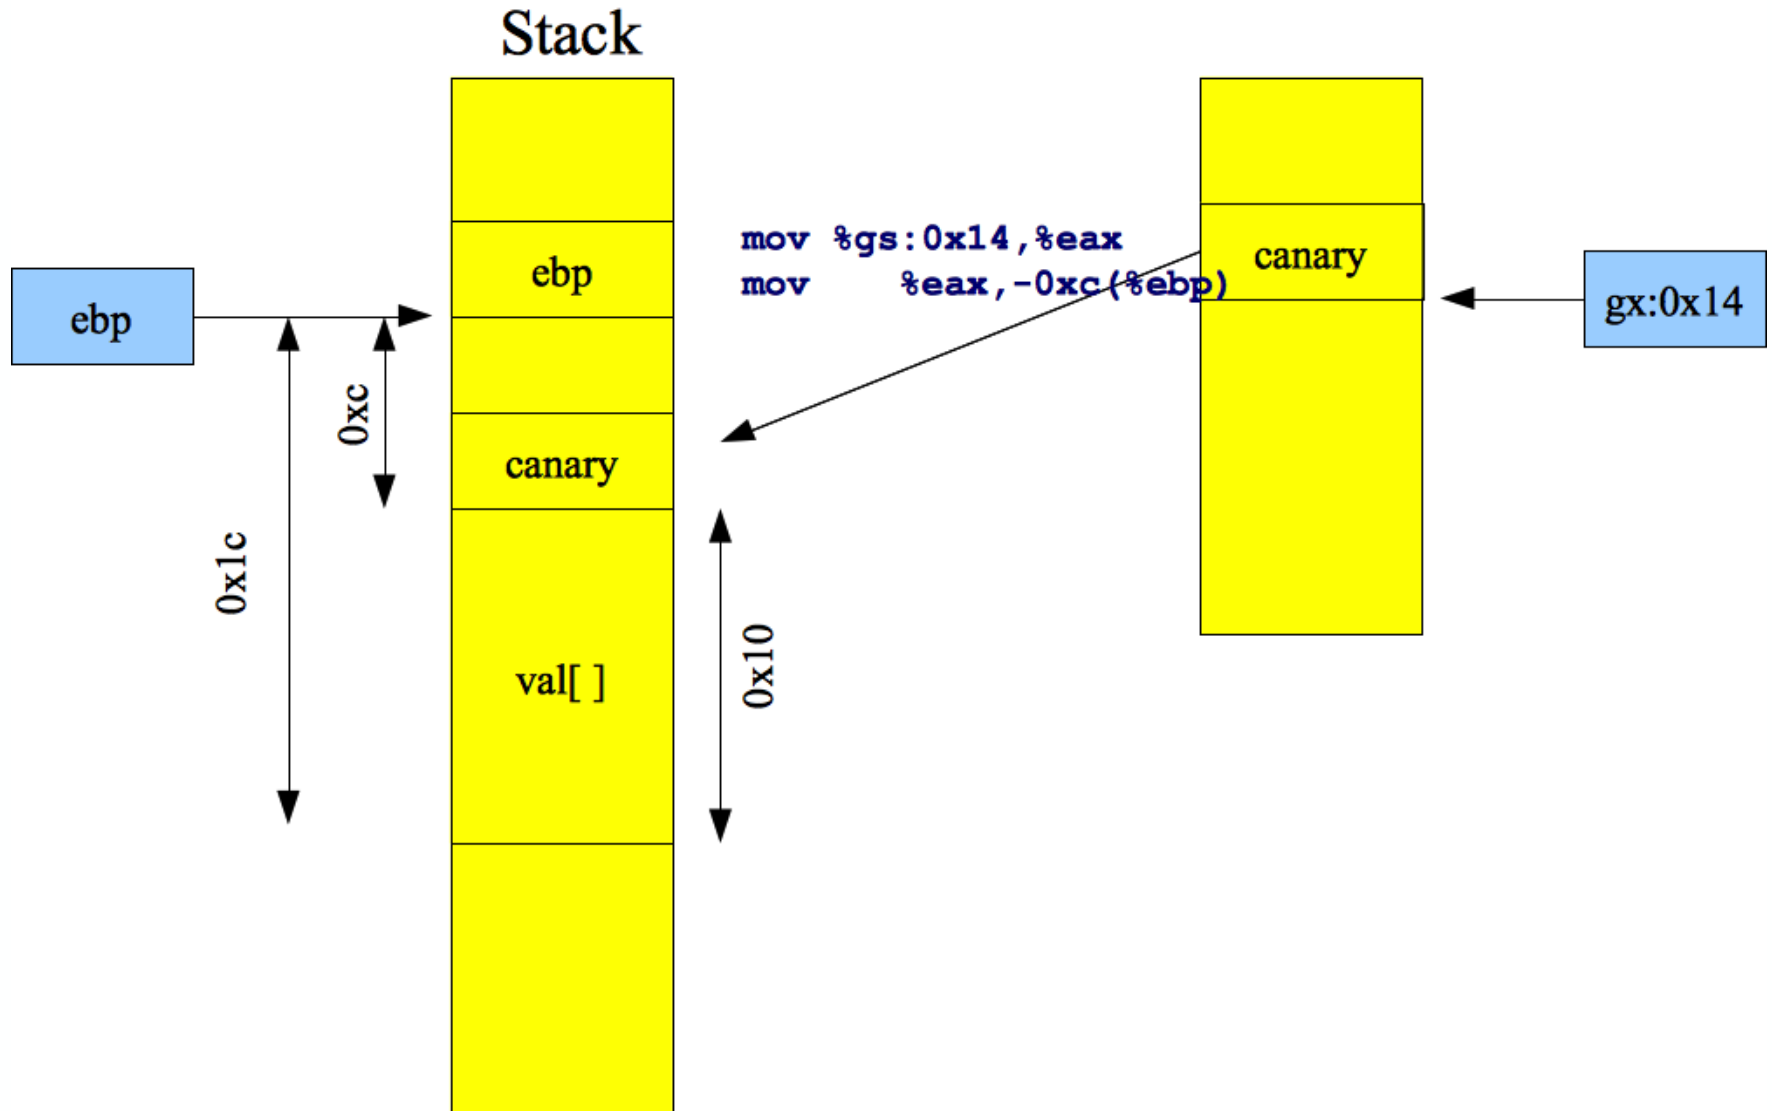
-Fstack-protection-all gcc option

Principe: localBufferOverflow assembler code (Intel code)

```
08048453 <_ZL19localBufferOverflowv>:
8048453:      55                push    %ebp
8048454:      89 e5             mov     %esp,%ebp
8048456:      83 ec 28          sub     $0x28,%esp
8048459:      65 a1 14 00 00 00  mov     %gs:0x14,%eax
804845f:      89 45 f4           mov     %eax,-0xc(%ebp)
8048462:      31 c0             xor     %eax,%eax
....
804848b:      8b 55 f4           mov     -0xc(%ebp),%edx
804848e:      65 33 15 14 00 00  xor     %gs:0x14,%edx
8048495:      74 05             je      804849c
8048497:      e8 54 fe ff ff    call    80482f0 <__stack_chk_fail@plt>
804849c:      c9                leave
```

If the edx register is not zero, the function `__stack_chk_fail@plt` is called

-Fstack-protection-all gcc option



-Fstack-protection-all gcc option, ARM uP

08048453 <testCanary>:

```
1049c:      e58de004      str     lr, [sp, #4]
104a0:      e28db004      add     fp, sp, #4
104a4:      e24dd010      sub     sp, sp, #16
104a8:      e3003f08      movw    r3, #3848      ; 0xf08
104ac:      e3403002      movt    r3, #2
104b0:      e5933000      ldr     r3, [r3]
104b4:      e50b3008      str     r3, [fp, #-8]
104b8:      e3a03000
...
104f8:      e3003f08      movw    r3, #3848      ; 0xf08
104fc:      e3403002      movt    r3, #2
10500:      e51b2008      ldr     r2, [fp, #-8]
10504:      e5933000      ldr     r3, [r3]
10508:      e1520003      cmp     r2, r3
1050c:      0a000000      beq     10514 <testCanary+0x7c>
10510:      ebffff8e      bl      10350 <__stack_chk_fail@plt>
10514:      e24bd004      sub     sp, fp, #4
10518:      e59db000      ldr     fp, [sp]
1051c:      e28dd004      add     sp, sp, #4
10520:      e49df004      pop     {pc}           ; (ldr pc, [sp],
```

If $r2 \neq r3$ --> `__stack_chk_fail` function is called

Randomize_va_space (CONFIG_COMPAT_BRK)

More information: [kernel source/Documentation/sysctl/kernel.txt](http://kernel.org/Documentation/sysctl/kernel.txt)

Make linux-xconfig: General Setup → [] Disable Heap randomization

This option can be used to select the **type of process address space randomization** (Called: Address space layout randomization (ASLR))

0 - Turn the process address space randomization off. This is the default for architectures that do not support this feature anyways, and kernels that are booted with the "**norandmaps**" parameter.

1 - Make the addresses of **mmap base, stack and VDSO (virtual dynamically linked shared objects) page randomized**. This, among other things, implies that shared libraries will be loaded to random addresses. This is the default if the CONFIG_COMPAT_BRK option is enabled.

2 - Additionally enable heap randomization. This is the default if **CONFIG_COMPAT_BRK** is disabled.

On the nanoPi:

```
cat /proc/sys/kernel/randomize_va_space           // show
echo value > /proc/sys/kernel/randomize_va_space   // modify
sysctl -w kernel.randomize_va_space=value          // modify value=0,1,2
```

Randomize_va_space (CONFIG_COMPAT_BRK)

Different commands on nanoPi:

Show the randomize_va_space value:

```
cat /proc/sys/kernel/randomize_va_space
```

Modify the randomize_va_space value:

```
echo 2 > /proc/sys/kernel/randomize_va_space
```

```
sysctl -w kernel.randomize_va_space=0 or 1 or 2
```

Hardening Linux:

```
echo 1 > /proc/sys/kernel/randomize_va_space
```

Better:

```
echo 2 > /proc/sys/kernel/randomize_va_space
```

```
sysctl -w kernel.randomize_va_space=2
```

Randomize address of kernel Image

(CONFIG_RANDOMIZE_BASE)

Randomizes virtual address at which the kernel image is loaded, as a security feature that deters exploit attempts relying on knowledge of the location of kernel internals.

Hardening Linux:

make linux-xconfig:

Kernel feature → [*] Randomize the address of the kernel image

Randomize SLAB Allocator

(CONFIG_SLAB_FREELIST_RANDOM, CONFIG_SLAB_FREELIST_HARDENED)

Make linux-xconfig: General Setup

[Choose SLAB allocator (SLAB) → [*] SLAB

[*] Allow slab caches to be merged

[*] SLAB freelist randomization

Slab allocation is a memory management mechanism intended for the efficient memory allocation of kernel objects. It randomizes the freelist order used on creating new pages. This security feature reduces the predictability of the kernel slab allocator against heap overflows

Write protect kernel text and module

(CONFIG_RODATA_FULL_DEFAULT_ENABLED)

Kernel code is in the text section and this section must be read only.

This prevents code or read-only data from being modified (inadvertently or intentionally) via another mapping of the same memory page.

Hardening Linux:

make linux-xconfig:

Kernel Feature → [*] Apply r/o permissions of VM areas also to their linear aliases

Optimize for performance or size

(CONFIG_CC_OPTIMIZE_FOR_PERFORMANCE, CC_OPTIMIZE_FOR_SIZE)

It is possible to optimize the kernel size (gcc -Os) or the performance (gcc -O2)

But for the current embedded systems, it is not necessary to optimize the kernel size

Make linux-xconfig: General Setup → Compiler optimization level

→ [*] Optimize for performance

Enable random number generator

dev.gentoo.org/~swift/docs/security_benchmarks/kernel.html

Enable the random number generator to provide a secure source of random numbers (which is important for cryptographic functions).

This can be accomplished by using the **CONFIG_ARCH_RANDOM** setting and next to select the random generator for the platform

Make linux-xconfig:

Device drivers → Character Devices → [*] Hardware Random Number Generator Core support

[*] Timer IOMEM HW Random Number Generator support

[*] Broadcom BCM2835/BCM63xx Random Number Generator support

Enable random number generator

- If the hardware used has no random generator, it is possible to use the HAVEGE [HArdware Volatile Entropy Gathering and Expansion] service (<https://www.irisa.fr/caps/projects/hipsor/>).
- It is an unpredictable random number generator is a practical approximation of a truly random number generator.

Important: `cd ~/workspace/nano/buildroot`

`make menuconfig`

→ Target Packages → Miscellaneous → haveged

...

On nanopi: The entropy can be read:

`cat /proc/sys/kernel/random/entropy-avail`

`sysctl kernel.random.entropy_avail`

Filter access to /dev/mem (STRICT_DEVMEM)

Do not allow all processes full access to all of the systems' memory through /dev/mem (which includes kernel memory and process memory).

By enabling CONFIG_STRICT_DEVMEM. Root user can only access memory regions expected for all legitimate common usage

Hardening Linux:

make linux-xconfig: Kernel Hacking → [*] Filter access to /dev/mem
→ [*] Filter I/O access to /dev/mem

Strip assembler-generated symbols during link

(STRIP_ASM_SYMS)

Strip internal assembler-generated symbols during a link, it is harder to reverse the code

Hardening Linux:

make linux-xconfig:

Kernel Hacking → Compile time check and compiler options

→ [] Compile the kernel with debug info

→ [*] Strip assembler-generated symbols during link

Restrict unprivileged access to the kernel syslog

(`SECURITY_DMESG_RESTRICT`)

Only root can access to the kernel system logs (through **dmesg**).

Hardening Linux:

make linux-xconfig: Security options → [*] Restrict unprivileged access to the kernel syslog

Kernel Memory Initialization (CONFIG_INIT_STACK_NONE, CONFIG_INIT_ON_ALLOC_DEFAULT_ON, CONFIG_INIT_ON_FREE_DEFAULT_ON)

This options enable initialization of stack variables and heap memory.

Hardening Linux:

make linux-xconfig: Security options → Kernel Hardening Options → Memory Initialization
→ Initialize Kernel Stack ... → () no automatic initialization (weakest) **(if possible)**
[*] Enable heap memory zeroing on allocation by default
[*] Enable heap memory zeroing on free by default

Check memory copies between kernel and

userspace (CONFIG_HARDENED_USERCOPY,
CONFIG_HARDENED_USERCOPY_FALLBACK
CONFIG_FORTIFY_SOURCE)

This option checks for obviously wrong memory regions when copying memory to/from the kernel (via `copy_to_user()` and `copy_from_user()` functions) by rejecting memory ranges that are larger than the specified heap object, span multiple separately allocated pages, are not on the process stack, or are part of the kernel text.

Hardening Linux:

make linux-xconfig: Security options

- [*] Harden memory copies between kernel and userspace
 - [*] Allow usercopy whitelist violations to fallback to object size
 - [*] Harden common str/mem functions against buffer overflows

File Systems

make linux-xconfig: File Systems →

(For each file system you use, make sure extended attribute support is enabled)

<*> Second extended fs support

[*] Ext2 extended attributes

[*] Ext2 POSIX Access Control Lists

[*] Ext2 Security Labels

<*> The extended 3 (ext3) file system

[*] Ext3 POSIX Access Control Lists

[*] Ext3 Security Labels

<*> The Extended 4 (ext4) file system

[*] Ext4 POSIX Access Control Lists

[*] Ext4 Security Labels

IPv4 protection

More information: <kernel source>/Documentation/networking/ip-sysctl.txt

https://dev.gentoo.org/~swift/docs/security_benchmarks/kernel.html#item-gt-sysctl-ipv4forward

Disable IPv4 Forwarding: **If the system is not as a router**, `ip_forward` must be disabled:

```
sysctl -w net.ipv4.ip_forward = 0
```

Enable Source Route Verification: **if the system is as a router**, `rp_filter` should be activated on all interfaces

```
sysctl -w net.ipv4.conf.INT.rp_filter = 1 //INT=eth0, lo, ...
```

Disable IP Source Routing: `accept_source_route` must be disable

```
sysctl -w net.ipv4.conf.INT.accept_source_route= 0 //  
INT=eth0, lo, ...
```

IPv4 protection

Disable ICMP Redirects: `accept_redirect` must be disable on all interfaces

```
sysctl -w net.ipv4.conf.INT.accept_redirects = 0 // INT=eth0, lo, ...
```

Ignore ICMP Echo Broadcasts: `icmp_echo_ignore_broadcasts` must be enable

```
sysctl -w net.ipv4.icmp_echo_ignore_broadcasts = 1
```

Ignore ICMP Bogus Error Responses: `icmp_ignore_bogus_error_responses` must be enable

```
sysctl -w net.ipv4.icmp_ignore_bogus_error_responses = 1
```

Enable Logging of Martians: `log_martians` should be enable on all interfaces

```
sysctl -w net.ipv4.conf.INT.log_martians=1 //INT=eth0, lo, ...
```


Enable TCP SYN Cookie: `tcp_syncookies` must be enable

```
sysctl -w net.ipv4.tcp_syncookies=1
```


Enable TCP SYN cookie protection

dev.gentoo.org/~swift/docs/security_benchmarks/kernel.html

Normal TCP connection



Half open TCP connection



Enable TCP SYN cookie protection

dev.gentoo.org/~swift/docs/security_benchmarks/kernel.html

Normal TCP/IP networking is open to an attack known as "SYN flooding".

This denial-of-service attack uses half open tcp connections.

The **CONFIG_SYN_COOKIES** can avoid this attack

make linux-xconfig: Networking support → Networking options → TCP IP Networking → IP TCP Sys Cookies [*]

Read syn cookies configuration

```
cat /proc/sys/net/ipv4/tcp_syncookies // 0= deactivated, 1=activated  
sysctl -n net.ipv4.tcp_syncookies
```

Hardening Linux: Activate sys cookies

```
echo 1 > cat /proc/sys/net/ipv4/tcp_syncookies  
sysctl net.ipv4.tcp_syncookies=1
```

Configure kernel parameters

`sysctl -p` command reads the file `/etc/sysctl.conf` and configure kernel parameters

Example:

```
cat /etc/sysctl.conf
kernel.randomize_va_space=2
net.ipv4.conf.lo.rp_filter = 1
net.ipv4.conf.eth0.rp_filter = 1
net.ipv4.conf.lo.accept_source_route= 0
net.ipv4.conf.eth0.accept_source_route= 0
net.ipv4.conf.lo.accept_redirects=0
net.ipv4.conf.eth0.accept_redirects=0
net.ipv4.icmp_echo_ignore_broadcasts = 1
net.ipv4.icmp_ignore_bogus_error_responses = 1
net.ipv4.conf.lo.log_martians=1
net.ipv4.conf.eth0.log_martians=1
net.ipv4.tcp_syncookies=1
```