

Valgrind

Jean-Roland Schuler

jeanrola.schuler.home.hefr.ch

Jean-roland.schuler@hefr.ch

References

[1]: <http://valgrind.org/>

[2]: <http://en.wikipedia.org/wiki/Valgrind>

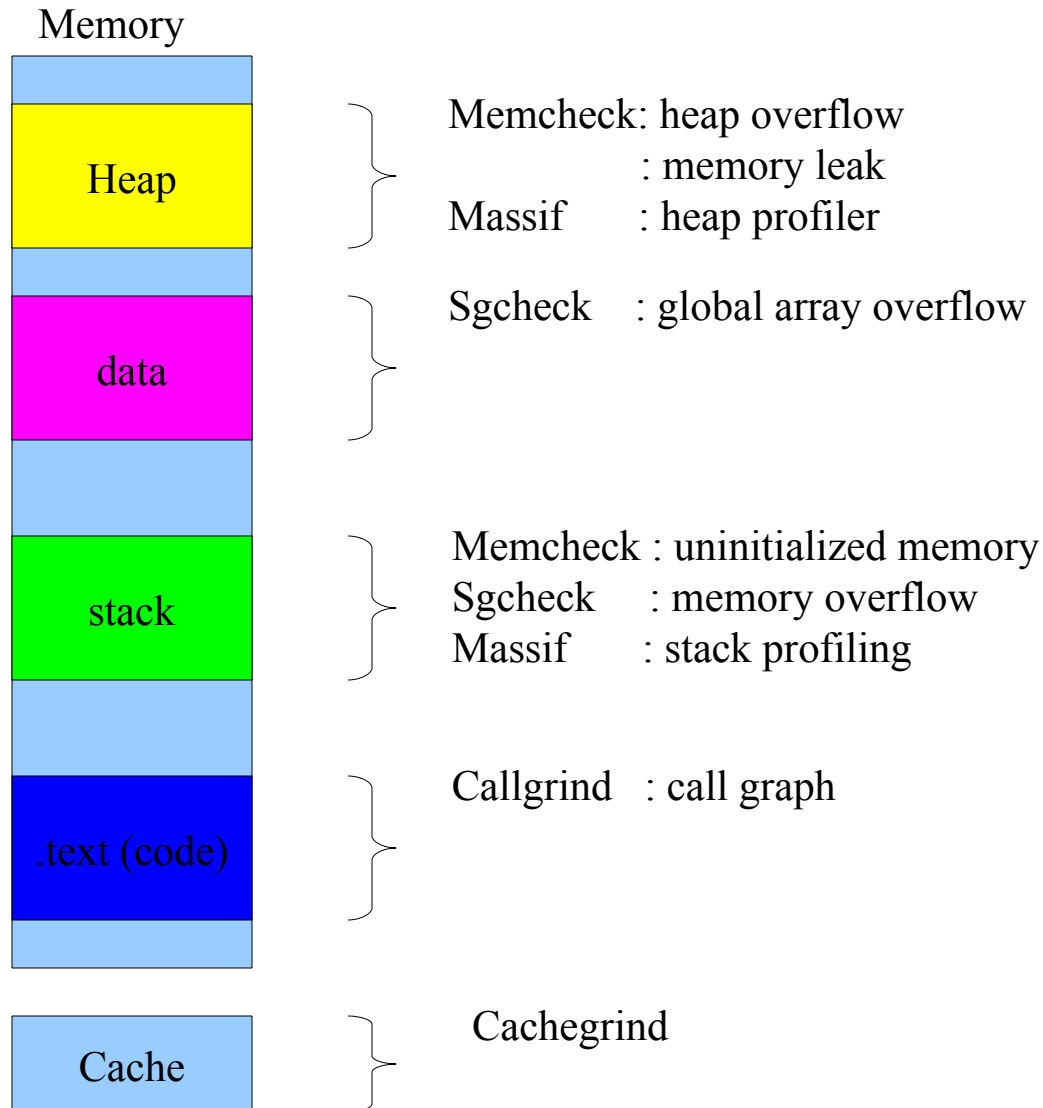
Valgrind ^{[1], [2]}

- Valgrind is an instrumentation framework for building dynamic analysis tools. There are Valgrind tools that can automatically detect many memory management and threading bugs, and profile your programs in detail. You can also use Valgrind to build new tools.
- Valgrind is in essence a virtual machine using just-in-time (JIT) compilation techniques, including dynamic recompilation. Nothing from the original program ever gets run directly on the host processor. Instead, Valgrind first translates the program into a temporary, simpler form called Intermediate Representation (IR), which is a processor-neutral. After the conversion, a tool (next page) is free to do whatever transformations it would like on the IR, before Valgrind translates the IR back into machine code and lets the host processor run it.
- It runs on the following platforms: X86/Linux, AMD64/Linux, ARM/Linux, PPC32/Linux, PPC64/Linux, S390X/Linux, MIPS/Linux, ARM/Android (2.3.x and later), X86/Android (4.0 and later), X86/Darwin and AMD64/Darwin, Mac OS X 10.12.

Valgrind: tools

1. **Memcheck** is a memory error detector. It helps you make your programs, particularly those written in C and C++, more correct.
2. **Cachegrind** is a cache and branch-prediction profiler. It helps you make your programs run faster.
3. **Callgrind** is a call-graph generating cache profiler. It has some overlap with Cachegrind, but also gathers some information that Cachegrind does not.
4. **Helgrind** is a thread error detector. It helps you make your multi-threaded programs more correct.
5. **DRD** is also a thread error detector. It is similar to Helgrind but uses different analysis techniques and so may find different problems.
6. **Massif** is a heap and stack profiler. It helps you make your programs use less memory.
7. **DHAT** is a tool for examining how programs use their heap allocations. It tracks the allocated blocks, and inspects every memory access to find which block, if any, it is to.

Valgrind summary



Memcheck

Memcheck is a memory error detector. It can detect the following problems that are common in C and C++ programs.

- Accessing memory you shouldn't, e.g. overrunning and under running heap blocks, and accessing memory after it has been freed.
- Using undefined values, i.e. values that have not been initialized, or that have been derived from other undefined values.
- Incorrect freeing of heap memory, such as double-freeing heap blocks, or mismatched use of *malloc-new* versus *free free/delete*
- Overlapping *src* and *dst* pointers in *memcpy* and related functions.
- Memory leaks.

Theses problems can be difficult to find, often remaining undetected for long periods, then causing occasional, difficult-to-diagnose crashes.

Memcheck: heap overflow

```
int mallocOverflow1(void)
{
    int i;
    int *p = malloc(sizeof(int) * 10);
    for (i = 0; i < 11; i++) {
        p[i] = i;
    }
    free(p);
    return 0;
}
```

Valgrind: Heap overflow

```
#gcc -Wall -g -o file file.c
```

```
#valgrind ./file          or  
#valgrind --tool=memcheck ./file
```

```
==6170== Command: ./code  
==6170== Invalid write of size 4  
==6170==    at 0x804858A: mallocOverflow1 (code.c:72)  
==6170==    by 0x80484FA: main (code.c:26)  
==6170== Address 0x4027050 is 0 bytes after a block of size 40  
alloc'd  
==6170==    at 0x4007E08: malloc (vg_replace_malloc.c:270)  
==6170==    by 0x804856B: mallocOverflow1 (code.c:70)  
==6170==    by 0x80484FA: main (file.c:26)  
==6170==  
==6170== All heap blocks were freed -- no leaks are possible  
==6170==  
==6170== For counts of detected and suppressed errors, rerun with: -v  
==6170== ERROR SUMMARY: 1 errors from 1 contexts (suppressed: 0 from  
0)
```


Memcheck: uninitialized memory

```
static int uninitializedMemory(void)
{
int i;
int a[10];
    for (i = 0; i < 9; i++)
        a[i] = i;

    for (i = 0; i < 10; i++)
    {
        printf("%d ", a[i]);
    }
    printf("\n");
    return 0;
}
```

Valgrind: uninitialized memory

```
#gcc -Wall -g -o file file.c
```

```
#valgrind ./file or
```

```
#valgrind -v ./file or
```

```
#valgrind --track-origins=yes -v ./file
```

```
#valgrind -s -v ./file // = --track-origins=yes
```

```
==3668== Command: ./file
```

```
==3668==
```

```
==3668== Use of uninitialised value of size 4
```

```
==3668== at 0x4DD4471B: _itoa_word (in /usr/lib/libc-2.15.so)
```

```
==3668== by 0x4DD491D8: vfprintf (in /usr/lib/libc-2.15.so)
```

```
==3668== by 0x4DD4EA6E: printf (in /usr/lib/libc-2.15.so)
```

```
==3668== by 0x80484FA: main (file.c:32)
```

```
...
```

```
==3668== ERROR SUMMARY: 25 errors from 7 contexts (suppressed: 0 from 0)
```

Memcheck: memory leak

```
int memoryLeak(void)
{
    int i;
    int *a;

    for (i=0; i < 10; i++)
    {
        a = malloc(sizeof(int) * 100);
    }
    free(a);
    return 0;
}
```

Memcheck: memory leak

```
#gcc -Wall -g -o file file.c
#valgrind ./file          or
#valgrind -v ./file       or
#valgrind --leak-check=full ./file
```

```
==3698== HEAP SUMMARY:
```

```
==3698==      in use at exit: 3,600 bytes in 9 blocks
```

```
==3698==    total heap usage: 10 allocs, 1 frees, 4,000 bytes
allocated
```

```
==3698==
```

```
==3698== LEAK SUMMARY:
```

```
==3698==      definitely lost: 3,600 bytes in 9 blocks
```

```
==3698==      indirectly lost: 0 bytes in 0 blocks
```

```
==3698==      possibly lost: 0 bytes in 0 blocks
```

```
==3698==      still reachable: 0 bytes in 0 blocks
```

```
==3698==      suppressed: 0 bytes in 0 blocks
```

```
==3698== Rerun with --leak-check=full to see details of leaked memory
=
```

Massif

- Massif is a heap profiler. It measures how much heap memory your program uses. This includes both the useful space, and the extra bytes allocated for book-keeping and alignment purposes. It can also measure the size of your program's stack(s), although it does not do so by default.

Massif: heap profiling

```
void g(void)
{
    malloc(4000);
}

void f(void)
{
    malloc(2000);
    g();
}

int main(void)
{
    int i;
    int* a[10];

    for (i = 0; i < 10; i++) {
        a[i] = malloc(1000);
    }
    f();
    g();
    for (i = 0; i < 10; i++) {
        free(a[i]);
    }
    return 0;
}
```

Massif: heap profiling

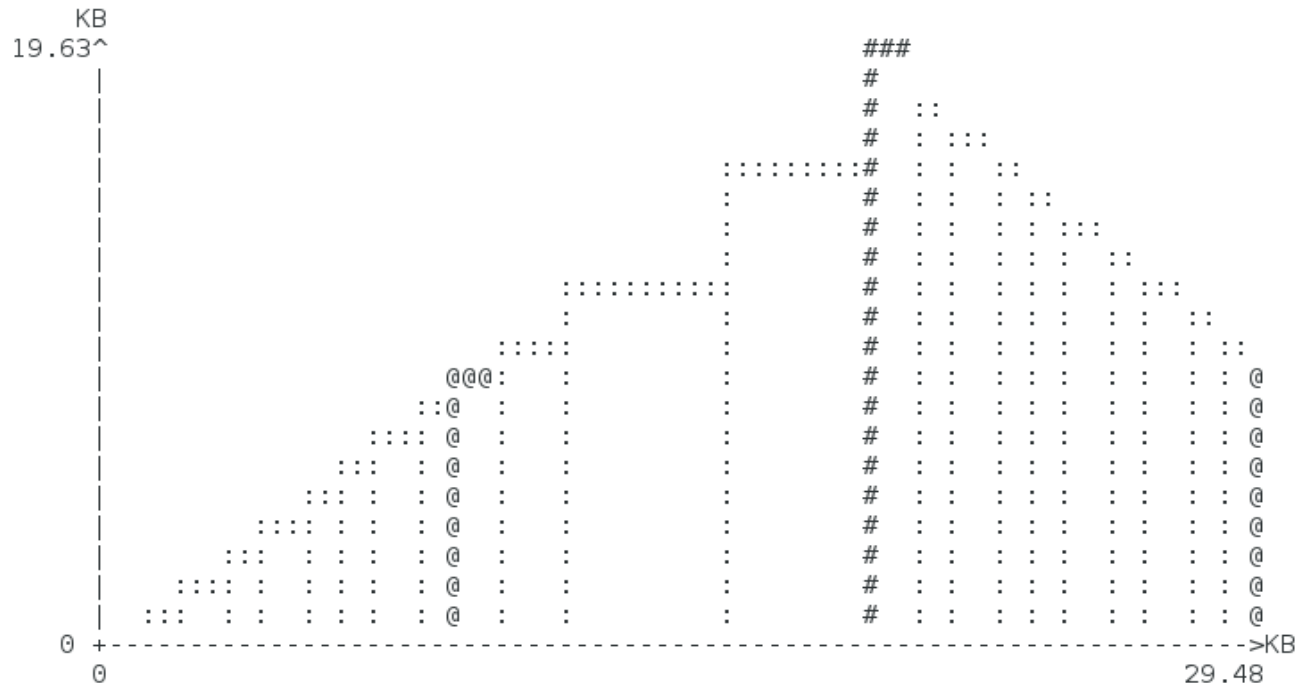
```
#gcc -Wall -g -o file file.c
```

```
#valgrind --tool=massif --time-unit=B ./file      (file massif.out.pid  
                                                    is generated)
```

```
#ms_print massif.out.PID
```

-time-unit: <i|ms|B> The time unit used for the profiling. i: Instructions executed (i), which is good for most cases; ms: Milliseconds, which is sometimes useful, B: Bytes allocated/deallocated, which is useful for very short-run programs

- Normal snapshots are represented in the graph by :
- Detailed snapshots are represented in the graph by @. There is a detailed snapshot every `-detailed_freq` option (default: 10)
- The peak snapshot is represented in the graph by: #



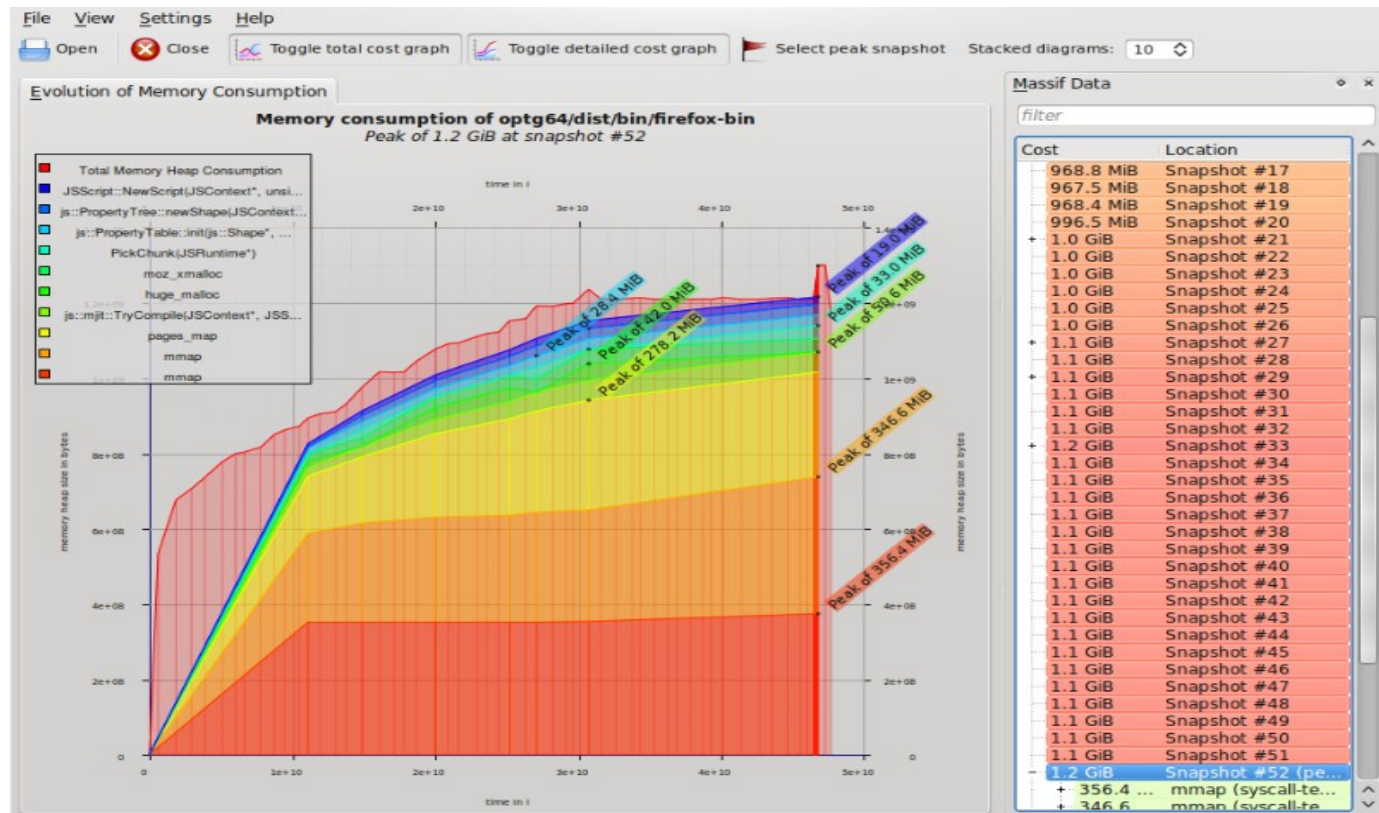
Massif: heap profiling

n	time (B)	total (B)	useful-heap (B)	extra-heap (B)	stacks (B)
0	0	0	0	0	0
1	1,008	1,008	1,000	8	0
2	2,016	2,016	2,000	16	0
3	3,024	3,024	3,000	24	0
4	4,032	4,032	4,000	32	0

n: Snapshot number
time(B): Time unit in byte
Total (B): Total memory allocated
Useful-heap (B): Memory used by the program
Extra-heap (B): Memory not used by the program (but used by the OS)
Stack (B): Memory stack profiling (by default off)

Massif: Analyze firefox

- <http://blog.mozilla.org/nnethercote/2010/12/09/memory-profiling-firefox-with-massif/>
- `valgrind --smc-check=all --trace-children=yes --tool=massif --pages-as-heap=yes --detailed-freq=1000000 optg64/dist/bin/firefox -P cad20 -no-remote`
- **Massif-visualizer** (kde-apps.org/content/show.php?content=122409)



Example: Memory lost

```
#gcc -Wall -g -o file file.c                (program p 18)
#valgrind ./file

==4164== HEAP SUMMARY:
==4164==      in use at exit: 10,000 bytes in 3 blocks
==4164==    total heap usage: 13 allocs, 10 frees, 20,000 bytes allocated
==4164==
==4164== LEAK SUMMARY:
==4164==    definitely lost: 10,000 bytes in 3 blocks
==4164==    indirectly lost: 0 bytes in 0 blocks
==4164==    possibly lost: 0 bytes in 0 blocks
==4164==    still reachable: 0 bytes in 0 blocks
==4164==    suppressed: 0 bytes in 0 blocks
==4164== Rerun with --leak-check=full to see details of leaked memory
```

Remark: With this command we know that memories are lost

Example: Memory lost

```
valgrind --leak-check=full ./file
```

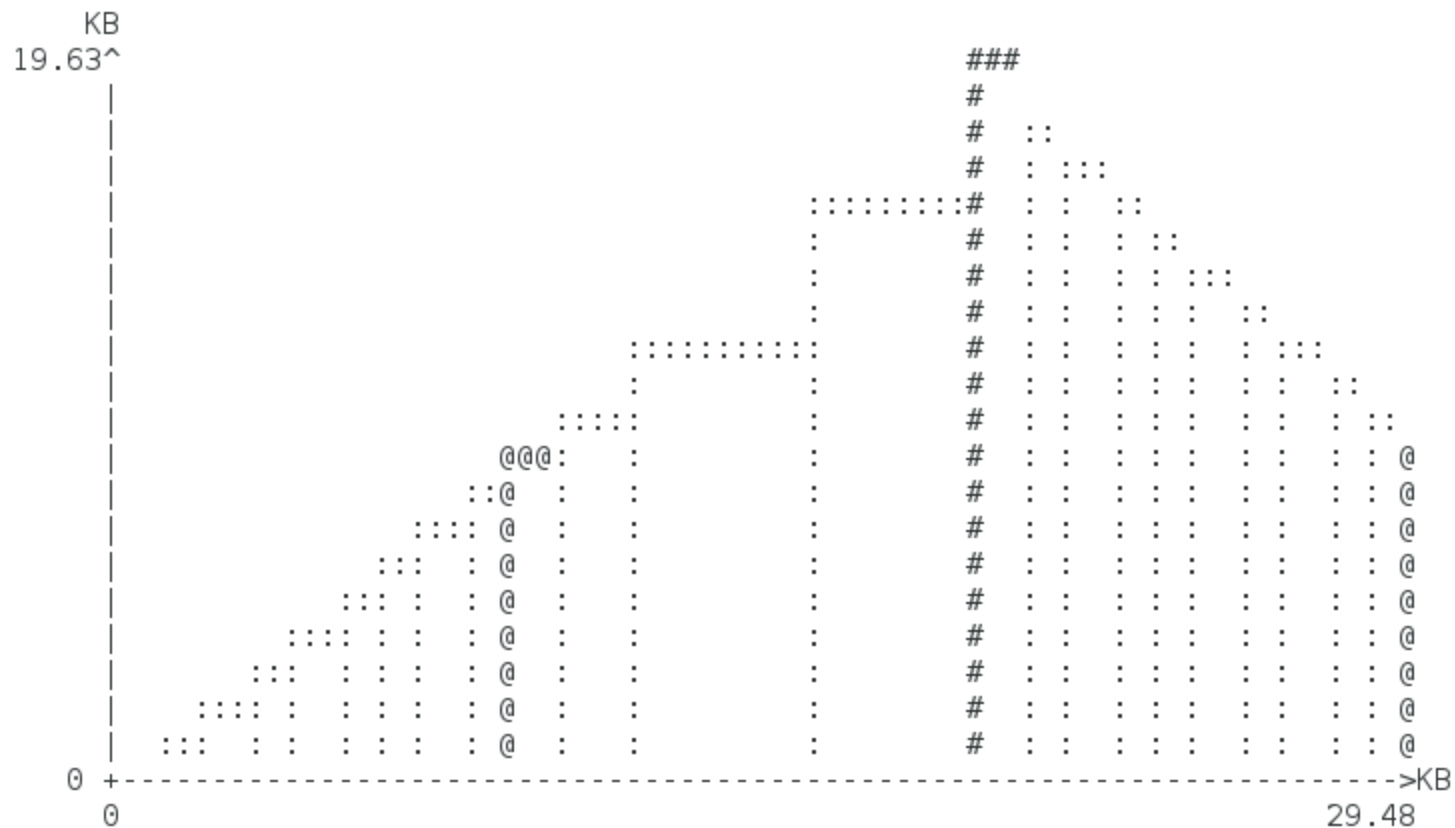
```
==4198== HEAP SUMMARY:
==4198==    in use at exit: 10,000 bytes in 3 blocks
==4198==    total heap usage: 13 allocs, 10 frees, 20,000 bytes allocated
==4198==
==4198== 2,000 bytes in 1 blocks are definitely lost in loss record 1 of 3
==4198==    at 0x4007E08: malloc (vg_replace_malloc.c:270)
==4198==    by 0x8048455: f (code1.c:9)
==4198==    by 0x8048496: main (code1.c:21)
==4198==
==4198== 4,000 bytes in 1 blocks are definitely lost in loss record 2 of 3
==4198==    at 0x4007E08: malloc (vg_replace_malloc.c:270)
==4198==    by 0x8048441: g (code1.c:4)
==4198==    by 0x804845A: f (code1.c:10)
==4198==    by 0x8048496: main (code1.c:21)
==4198==
==4198== 4,000 bytes in 1 blocks are definitely lost in loss record 3 of 3
==4198==    at 0x4007E08: malloc (vg_replace_malloc.c:270)
==4198==    by 0x8048441: g (code1.c:4)
==4198==    by 0x804849B: main (code1.c:22)
==4198==
==4198== LEAK SUMMARY:
==4198==    definitely lost: 10,000 bytes in 3 blocks
==4198==    indirectly lost: 0 bytes in 0 blocks
==4198==    possibly lost: 0 bytes in 0 blocks
==4198==    still reachable: 0 bytes in 0 blocks
==4198==    suppressed: 0 bytes in 0 blocks
```

Remark: This command shows where the memories are lost

Example:Memory lost

```
#valgrind --tool=massif --time-unit=B ./file
#ms_printf massif.out.pid
```

Remark: this graph tells the evolution of the heap memory



Example: Memory lost

n	time (B)	total (B)	useful-heap (B)	extra-heap (B)	stacks (B)
0	0	0	0	0	0
1	1,008	1,008	1,000	8	0
2	2,016	2,016	2,000	16	0
3	3,024	3,024	3,000	24	0
4	4,032	4,032	4,000	32	0
5	5,040	5,040	5,000	40	0
6	6,048	6,048	6,000	48	0
7	7,056	7,056	7,000	56	0
8	8,064	8,064	8,000	64	0
9	9,072	9,072	9,000	72	0

99.21% (9,000B) (heap allocation functions) malloc/new/new[], --alloc-fns, etc.
 ->99.21% (9,000B) 0x804847A: main (**code1.c:19**)

1st detailed snapshot,
indicate where the
memory is allocated

n	time (B)	total (B)	useful-heap (B)	extra-heap (B)	stacks (B)
10	10,080	10,080	10,000	80	0
11	12,088	12,088	12,000	88	0
12	16,096	16,096	16,000	96	0
13	20,104	20,104	20,000	104	0
14	20,104	20,104	20,000	104	0

99.48% (20,000B) (heap allocation functions) malloc/new/new[], --alloc-fns, etc.
 ->49.74% (10,000B) 0x804847A: main (**code1.c:19**)
 |
 ->39.79% (8,000B) 0x8048440: g (**code1.c:4**)
 | ->19.90% (4,000B) 0x8048459: f (code1.c:10)
 | | ->19.90% (4,000B) 0x8048495: main (code1.c:21)
 | |
 | ->19.90% (4,000B) 0x804849A: main (code1.c:22)
 |
 ->09.95% (2,000B) 0x8048454: f (**code1.c:9**)
 ->09.95% (2,000B) 0x8048495: main (code1.c:21)

Peak snapshot, with
the detail

Example: Memory lost

n	time (B)	total (B)	useful-heap (B)	extra-heap (B)	stacks (B)
15	21,112	19,096	19,000	96	0
16	22,120	18,088	18,000	88	0
17	23,128	17,080	17,000	80	0
18	24,136	16,072	16,000	72	0
19	25,144	15,064	15,000	64	0
20	26,152	14,056	14,000	56	0
21	27,160	13,048	13,000	48	0
22	28,168	12,040	12,000	40	0
23	29,176	11,032	11,000	32	0
24	30,184	10,024	10,000	24	0

99.76% (10,000B) (heap allocation functions) malloc/new/new[], --alloc-fns, etc.
->79.81% (8,000B) 0x8048440: g (**code1.c:4**)
| ->39.90% (4,000B) 0x8048459: f (code1.c:10)
| | ->39.90% (4,000B) 0x8048495: main (code1.c:21)
| |
| ->39.90% (4,000B) 0x804849A: main (code1.c:22)
|
->19.95% (2,000B) 0x8048454: f (**code1.c:9**)
| ->19.95% (2,000B) 0x8048495: main (code1.c:21)
|
->00.00% (0B) in 1+ places, all below ms_print's threshold (01.00%)

Lost memories, and where
the memory is allocated

Massif: example, stack profiling

```
void f(void)
{
    char buffer [1000];

}

int main(void)
{
    char buffer [100];

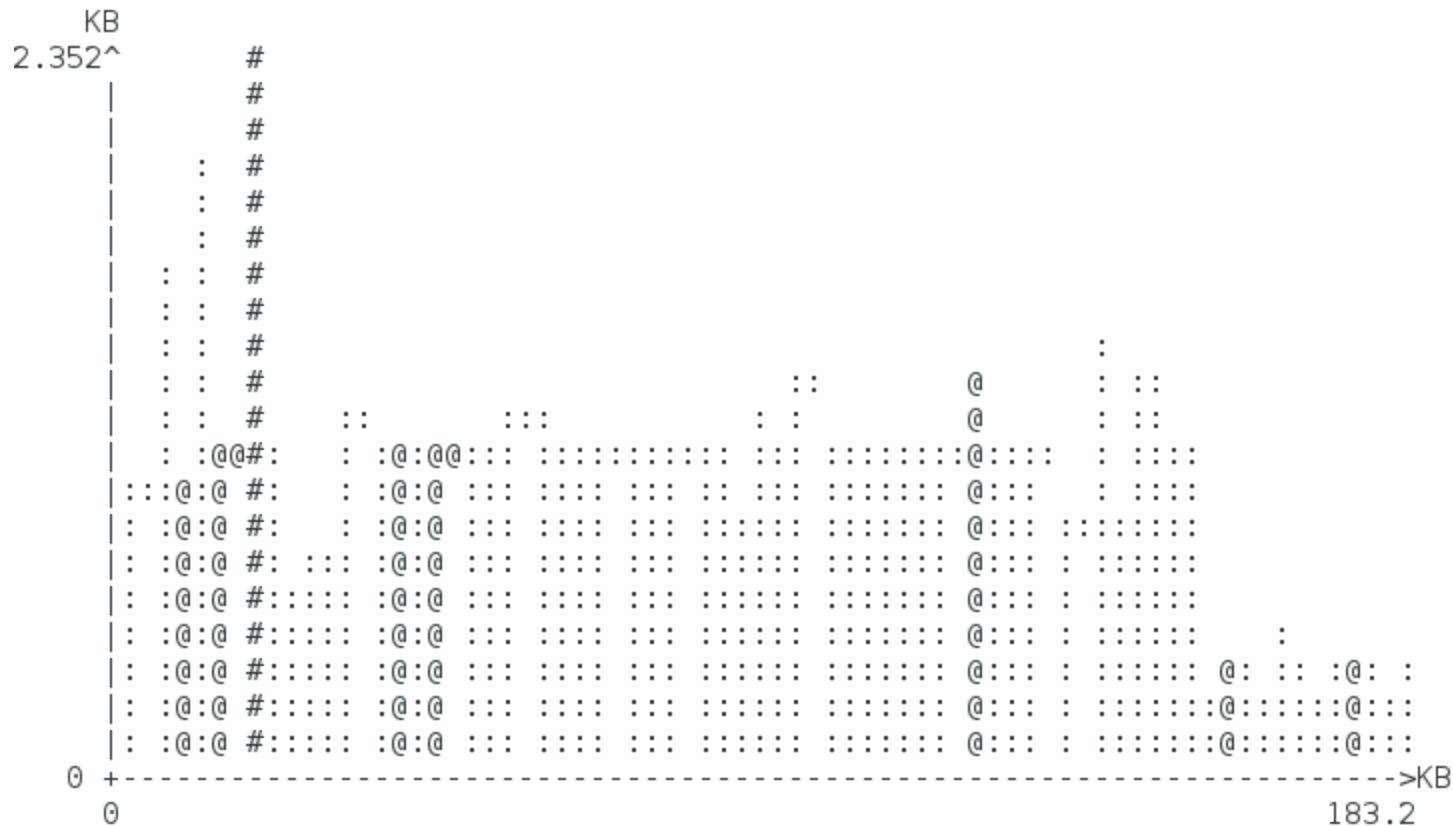
    f();
    return 0;
}
```

Massif: example, stack profiling

```
#gcc -Wall -g -o file file.c
```

```
#valgrind --tool=massif --time-unit=B --stacks=yes ./file
```

```
#ms_print massif.out.PID
```



Callgrind

- Callgrind: a call-graph generating cache and branch prediction profiler
- Callgrind is a profiling tool that records the call history among functions in a program's run as a call-graph.
- By default, the collected data consists of:
 - the number of instructions executed,
 - their relationship to source lines,
 - the caller/callee relationship between functions,
 - and the numbers of such calls.
- Optionally, cache simulation and/or branch prediction (similar to Cachegrind) can produce further information about the runtime behavior of an application.

Callgrind: example

```
// test.c  
//gcc -Wall -o test test.c test_new.c
```

```
#include<stdio.h>
```

```
void new_func1(void);
```

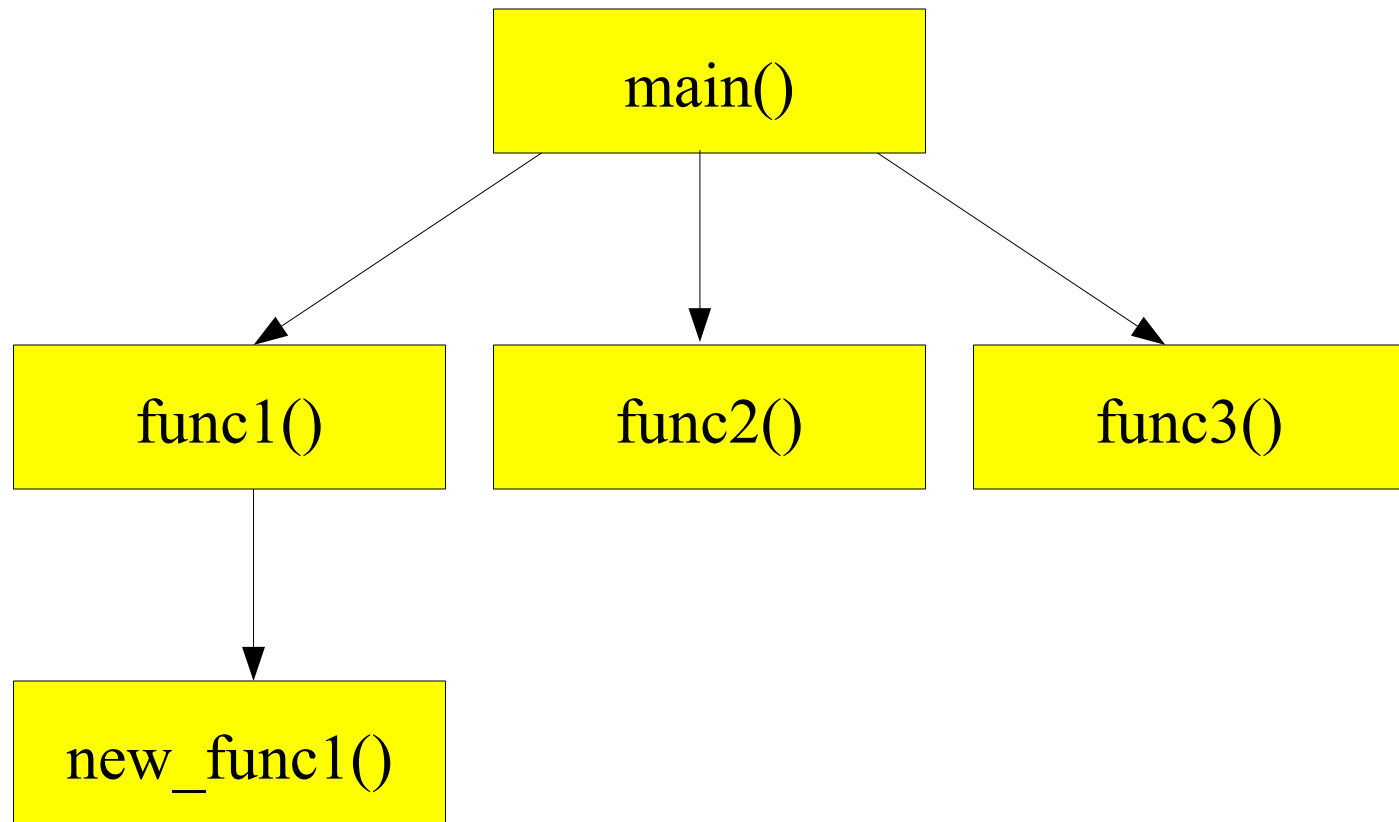
```
void func1(void)  
{  
    printf("\n Inside func1 \n");  
    int i = 0;  
  
    for(;i<0xffff;i++);  
    new_func1();  
    return;  
}
```

```
static void func2(void)  
{  
    printf("\n Inside func2 \n");  
    int i = 0;  
  
    for(;i<0xffaa;i++);  
    return;  
}
```

```
static void func3(void)  
{  
    printf("\n Inside func2 \n");  
    return;  
}  
  
int main(void)  
{  
    printf("\n Inside main() \n");  
    int i = 0;  
  
    for(;i<0xffff;i++);  
    func1();  
    func2();  
    func3();  
  
    return 0;  
}
```

```
//test_new.c  
void new_func1(void) {  
    printf("\n Inside new_func1() \n");  
    int i = 0;  
  
    for(;i<0xffee;i++);  
  
    return;  
}
```

Callgrind: example



Callgrind: example

```
#gcc -Wall -g -o test test.c test_new.c
```

```
#valgrind --tool=callgrind ./test
```

- This command generates a file: *callgrind_out.PID*

```
#kcachegrind
```

Callgrind: example

./callgrind.out.1840 [./test]

File View Go Settings Help

Open Back Forward Up Relative Cycle Detection Relative to Parent Shorten Templates Instruction Fetch

Flat Profile

Search: ELF Object

Self	ELF Object
84.61	test
13.13	ld-2.15.so
2.25	libc-2.15.so
0.01	vgpreload_core-x86-linux.so
0.00	(unknown)

Incl.	Self	Called	Function	Location
86.27	0.00	1	Ox08048300	test
85.84	2.04	1	main	test: test.c
55.74	32.66	1	func1	test: test.c
27.89	32.66	1	new_func1	test: test_new.c
27.82	32.62	1	func2	test: test.c
0.22	0.00	1	func3	test: test.c
0.01	0.01	1	__libc_csu_init	test
0.00	0.00	1	Ox080483a0	test
0.00	0.00	1	Ox080483d0	test
0.00	0.00	3	Ox08048324	test
0.00	0.00	1	Ox08048330	test

main

Types Callers All Callers Callee Map Source Code

```
# Ir CEst Source (/home/vmshared/work/Security/master/cours/coursSoftwareSecurity/prog
...
for(i<0xffff;i++);
new_func1();
```

```
graph TD
    main[main 100.00%] --> func2[func2 32.40%]
    main --> func1[func1 64.93%]
    func1 --> new_func1[new_func1 32.49%]
```

Parts Callees Call Graph All Callees Caller Map Machine Code

Cachegrind

- Cachegrind: a cache and branch-prediction profiler
- Cachegrind simulates how your program interacts with a machine's cache hierarchy and (optionally) branch predictor.
- It simulates a machine with independent first-level instruction and data caches (I1 and D1), backed by a unified second-level cache (L2). This exactly matches the configuration of many modern machines.
- However, some modern machines have three levels of cache. For these machines (in the cases where Cachegrind can auto-detect the cache configuration) Cachegrind simulates the first-level and third-level caches. The reason for this choice is that the L3 cache has the most influence on runtime, as it masks accesses to main memory.
- Therefore, Cachegrind always refers to the I1, D1 and LL (last-level) caches.

Cachegrind

- Cachegrind gathers the following statistics (abbreviations used for each statistic is given in parentheses):
- **I cache reads** (**Ir**, which equals the number of instructions executed), I1 cache read misses (**I1mr**) and LL cache instruction read misses (**ILmr**).
- **D cache reads** (**Dr**, which equals the number of memory reads), D1 cache read misses (**D1mr**), and LL cache data read misses (**DLmr**).
- **D cache writes** (**Dw**, which equals the number of memory writes), D1 cache write misses (**D1mw**), and LL cache data write misses (**DLmw**).
- Conditional branches executed (**Bc**) and conditional branches mispredicted (**Bcm**).
- Indirect branches executed (**Bi**) and indirect branches mispredicted (**Bim**).
- Note that **D1** total accesses is given by **D1mr + D1mw**, and that **LL** total accesses is given by **ILmr + DLmr + DLmw**.

Cachegrind: example

```
//gcc -Wall -o test test.c test_new.c

#include<stdio.h>

#define VAL1      5000
#define VAL2      100

static int x[VAL1][VAL2];

static int badFunction(void) {
    int      i, j;

    for (j=0; j<VAL2; j++) {
        for (i=0; i<VAL1; i++)
            x[i][j] = x[i][j]*2;
    }
    return 0;
}

static int goodFunction(void) {
    int      i, j;

    for (i=0; i<VAL1; i++)
    {
        for (j=0; j<VAL2; j++)
            x[i][j] = x[i][j]*2;
    }
    return 0;
}
```

```
int main (void)
{

    badFunction();
    goodFunction();
    return (0);
}
```

X[0][0]
X[0][1]
..
..
X[1][0]
X[1][1]
..
..



x[4999][99]

Cachegrind: example

```
#gcc -Wall -g -o test test.c
```

```
#valgrind --tool=cachegrind ./test
```

- This command generates a file: *cachegrind_out.PID*

```
#cg_annotate -auto=yes cachegrind_out.PID
```

Cachegrind: example

```
I1 cache:      32768 B, 64 B, 8-way associative
D1 cache:      32768 B, 64 B, 8-way associative
LL cache:      2097152 B, 64 B, 8-way associative
Command:       ./l_test
Data file:     cachegrind.out.3489
Events recorded: Ir I1mr I1Lmr Dr D1mr DLmr Dw D1mw DLmw
Events shown:   Ir I1mr I1Lmr Dr D1mr DLmr Dw D1mw DLmw
Event sort order: Ir I1mr I1Lmr Dr D1mr DLmr Dw D1mw DLmw
Thresholds:     0.1 100 100 100 100 100 100 100 100
Include dirs:
User annotated:
Auto-annotation: on
```

Cachegrind: example

	Ir	Ilmr	ILmr	Dr	Dlmr	DLmr	Dw	Dlmw	DLmw	
.
.	static int badFunction(void)
3	0	0	0	0	0	0	1	0	0	{
.	int i, j;
.	
304	1	1	1	201	0	0	1	0	0	for (j=0; j<VAL2; j++)
.	{
1,500,400	0	0	0	1,000,100	0	0	100	0	0	for (i=0; i<VAL1; i++)
5,500,000	0	0	0	2,500,000	500,000	31,251	500,000	0	0	x[i][j] = x[i][j]*2;
.	}
1	0	0	0	0	0	0	0	0	0	return 0;
2	0	0	0	2	0	0	0	0	0	}
.	
.	static int goodFunction(void)
3	0	0	0	0	0	0	1	0	0	{
.	int i, j;
.	
15,004	1	1	1	10,001	0	0	1	0	0	for (i=0; i<VAL1; i++)
.	{
1,520,000	0	0	0	1,005,000	0	0	5,000	0	0	for (j=0; j<VAL2; j++)
5,500,000	0	0	0	2,500,000	31,251	0	500,000	0	0	x[i][j] = x[i][j]*2;
.	}
1	0	0	0	0	0	0	0	0	0	return 0;
2	0	0	0	2	0	0	0	0	0	}
.	